

Number Representation

Mainak Chaudhuri
Indian Institute of Technology Kanpur

Sketch

- Unsigned and signed integers
- Overflow in integer operations
- Floating-point numbers
- Overflow in floating-point operations
- Underflow in floating-point operations
- Rounding modes

Unsigned and signed integers

- Numbers in a computer (or a processor) are represented in binary
 - Digital designs are realized using devices such as transistors which can be made to operate at two discrete voltage levels representing 0 and 1 conveniently
 - Number of bits devoted to represent a number is fixed by the design and cannot be extended after the design is finished
 - Puts an upper bound on the value of the largest representable number in a computer
 - If a computation generates a number that cannot fit within these bits, an overflow is said to have occurred

Unsigned and signed integers

- Non-negative integers (often called unsigned integers) can be represented using standard binary
- Negative integers require the sign to be encoded appropriately
 - Four possibilities have been tried: sign magnitude, two's complement, one's complement, biased
 - Sign magnitude is the simplest: reserve the most significant bit to represent the sign of the integer
 - Ambiguous representation of zero (00...0 and 100...0)
 - Addition requires extra logic to set the result's sign

Representing negative integers

- One's complement representation
 - An integer x and its negative add up to $2^n - 1$
 - $x + (-x) = 2^n - 1$ or $(-x) = 2^n - 1 - x$
 - Just bitwise inversion of x is $-x$
 - Ambiguous representation of zero
 - $00\dots 0$ and $11\dots 1$
 - How easy is addition or subtraction?
 - Various cases have different treatments
 - Not as easy as two's complement addition/subtraction
 - Balanced ranges on positive and negative sides
 - $-2^{n-1} + 1$ to $2^{n-1} - 1$

Representing negative integers

- Biased representation
 - Representation of an integer x is $x + \text{bias}$
 - $x + \text{bias}$ is represented as a regular binary number with no sign bit
 - bias is a constant fixed by the representation
 - The goal is to avoid negative numbers all together i.e., $x + \text{bias} \geq 0$
 - Smallest representable number is $-\text{bias}$
 - A typical bias is 2^{n-1} for n -bit representation
 - An integer x is represented as $x + 2^{n-1}$ using unsigned binary representation
 - Imbalanced range on positive and negative sides
 - $-\text{bias}$ to $2^n - 1 - \text{bias}$
 - -2^{n-1} to $2^{n-1} - 1$ when bias is 2^{n-1}

Representing negative integers

- Use of two's complement representation
 - All computers use this representation
 - The following C code can be used to verify whether a computer is using two's complement representation for 32-bit integers ("int" type)


```
int x = -10; // Any value can be used
unsigned bits = *((unsigned*)&x);
printf ("bits = %#x\n", bits);
```
 - Similarly, 64-bit integers ("long long" type) can be verified


```
long long x = -10;
unsigned long long bits = *((unsigned long long*)&x);
printf ("bits = %#llx\n", bits);
```

Overflow in integer operations

- Overflow occurs if the result of a computation cannot fit within the given number of bits
 - In usual binary representation, if there is a carry-out in the MSB position of an addition, overflow is said to occur
 - In two's complement representation, overflow occurs if and only if the carry in to the MSB is different from carry out from the MSB
 - Proof?
 - Let the carry in to the MSB be c_{n-1} and carry out from the MSB be c_n when adding two n -bit numbers

Overflow in integer operations

- Overflow in two's complement addition
 - Suppose that we are adding A ($a_{n-1}a_{n-2}...a_0$) and B ($b_{n-1}b_{n-2}...b_0$) both in two's complement representation and let the sum be S ($s_{n-1}s_{n-2}...s_0$)
 - Case I: $c_{n-1}=0, c_n=1$
 - Therefore, $a_{n-1}=b_{n-1}=1$ and $s_{n-1}=0$
 - This is an overflow condition because adding two negative numbers cannot yield a positive number
 - Case II: $c_{n-1}=1, c_n=0$
 - Therefore, $a_{n-1}=b_{n-1}=0$ and $s_{n-1}=1$
 - This is an overflow condition because adding two positive numbers cannot yield a negative number

Overflow in integer operations

- Overflow in two's complement addition
 - Consider adding A ($a_{n-1}a_{n-2}...a_0$) and B ($b_{n-1}b_{n-2}...b_0$) both in two's complement representation and let the sum be S ($s_{n-1}s_{n-2}...s_0$)
 - Case III: $c_{n-1}=0, c_n=0$
 - Therefore, at most one of a_{n-1} and b_{n-1} is 1
 - Case IIIA: $a_{n-1}=b_{n-1}=0$
 - Since c_{n-1} is 0, the magnitudes of two positive numbers are added without any overflow
 - Therefore, there is no overflow
 - Case IIIB: $a_{n-1}=0, b_{n-1}=1$ ($a_{n-1}=1, b_{n-1}=0$ is similar)
 - Therefore, $A+B = -2^{n-1} + \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$
 - Since c_{n-1} is 0, there is no overflow in adding lower $n-1$ bits of A and B; hence, $s_{n-2}s_{n-3}...s_0 = \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$
 - So, $A+B = -2^{n-1} + s_{n-2}s_{n-3}...s_0 = 1s_{n-2}s_{n-3}...s_0$ (no overflow)

Overflow in integer operations

- Overflow in two's complement addition
 - Consider adding A ($a_{n-1}a_{n-2}...a_0$) and B ($b_{n-1}b_{n-2}...b_0$) both in two's complement representation and let the sum be S ($s_{n-1}s_{n-2}...s_0$)
 - Case IV: $c_{n-1}=1, c_n=1$
 - Therefore, at most one of a_{n-1} and b_{n-1} is 0
 - Case IVA: $a_{n-1}=1, b_{n-1}=0$ ($a_{n-1}=0, b_{n-1}=1$ is similar)
 - Therefore, $A+B = -2^{n-1} + \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$
 - Since c_{n-1} is 1, the result of adding lower $n-1$ bits of A and B is $1s_{n-2}s_{n-3}...s_0 = \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$, where $1s_{n-2}s_{n-3}...s_0$ is treated as a positive value i.e., $2^{n-1} + \sum_{i=0}^{n-2} s_i 2^i$
 - So, $A+B = -2^{n-1} + 1s_{n-2}s_{n-3}...s_0 = -2^{n-1} + 2^{n-1} + \sum_{i=0}^{n-2} s_i 2^i = 0s_{n-2}s_{n-3}...s_0$
 - Since in this case, $a_{n-1}+b_{n-1}+c_{n-1} = c_n=0$, correct result is obtained by ignoring c_n (hence, no overflow)

Overflow in integer operations

- Overflow in two's complement addition
 - Consider adding A ($a_{n-1}a_{n-2}...a_0$) and B ($b_{n-1}b_{n-2}...b_0$) both in two's complement representation and let the sum be S ($s_{n-1}s_{n-2}...s_0$)
 - Case IV: $c_{n-1}=1, c_n=1$
 - Therefore, at most one of a_{n-1} and b_{n-1} is 0
 - Case IVB: $a_{n-1}=1, b_{n-1}=1$
 - Therefore, $A+B = -2^n + \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$
 - Since c_{n-1} is 1, the result of adding lower $n-1$ bits of A and B is $1s_{n-2}s_{n-3}...s_0 = \sum_{i=0}^{n-2} (a_i + b_i) \cdot 2^i$, where $1s_{n-2}s_{n-3}...s_0$ is treated as a positive value i.e., $2^{n-1} + \sum_{i=0}^{n-2} s_i 2^i$
 - So, $A+B = -2^n + 1s_{n-2}s_{n-3}...s_0 = -2^n + 2^{n-1} + \sum_{i=0}^{n-2} s_i 2^i = -2^{n-1} + \sum_{i=0}^{n-2} s_i 2^i = 1s_{n-2}s_{n-3}...s_0$ in two's complement
 - Since in this case, $a_{n-1}+b_{n-1}+c_{n-1} = c_n=1$, correct result is obtained by ignoring c_n (hence, no overflow)

Floating-point numbers

- Real numbers are represented using scientific notation
 - $a.b \times 10^c$
 - Said to be normalized if there is only one non-zero digit to the left of the decimal point
 - $1 \leq a \leq 9$
 - Always possible to normalize a real number
- Digital computers represent real numbers as floating-point binary numbers using normalized scientific notation
 - $1.b \times 2^c$ where b and c are binary numbers
 - Binary point as opposed to decimal point

Floating-point numbers

- Normalized scientific representation
 - Floating-point because the binary point can be moved by adjusting the exponent
- How to convert a decimal real number to normalized scientific binary notation?
 - Procedure for integer part: continued division
 - Procedure for fraction part: continued multiplication
- A representation of normalized floating-point number fixes the number of bits needed for b and c , given a fixed total number of bits
 - Trade-off between range and precision
 - If b has more bits then c has less

Floating-point numbers

- Normalized scientific representation has three advantages
 - Simplifies exchange of floating-point data between computers due to a standard representation
 - Simplifies floating-point arithmetic algorithms because all operands are in the same format
 - Compacts the representation by discarding leading zeros on the left hand side
 - 0.0000000010101 is same as 1.0101×2^{-9} and it is enough to store $b=0101$ and $c=-9$
 - b is referred to as the fraction (or mantissa) and c is referred to as the exponent

Floating-point numbers

- Normalized scientific representation
 - IEEE 754 standard dictates the number of bits reserved for mantissa and exponent
 - It is a standard format for representing floating-point numbers
 - All computers are expected to follow this format
 - A single-precision floating-point number is represented using 32 bits
 - Sign magnitude representation
 - MSB is sign bit
 - Least significant 23 bits represent the mantissa
 - 24 bits of (1 + mantissa) represent the significand
 - Middle eight bits represent the exponent

Floating-point numbers

- IEEE 754 single-precision format
 - Exponent field needs to encode both positive and negative exponents
 - Necessary to choose an encoding such that the entire 31-bit magnitude field is monotonic in the value of the floating-point number
 - A larger magnitude floating-point number should have a larger 31-bit magnitude field compared to a smaller magnitude floating-point number
 - Helps to sort numbers quickly by magnitude
 - Mantissa is already monotonic in the value of the fraction
 - Cannot use two's complement encoding for exponent
 - Negative exponents will be represented by large numbers

Floating-point numbers

- IEEE 754 single-precision format
 - Biased encoding of exponent with bias set to 127
 - Encoded exponent = 127 + actual exponent
 - Actual exponent is allowed to range from -126 to 127 i.e., the encoded exponent can range from 00000001 to 11111110
 - Encoded exponent cannot be 00000000 or 11111111 because these encodings are reserved to represent some special numbers
 - Notice that larger encoded exponent now represents larger actual exponent because biased encoding preserves monotonicity
 - Two's complement encoding does not preserve monotonicity

Floating-point numbers

- IEEE 754 single-precision format
 - Largest non-negative number

0	11111110	11111111111111111111111111111111
---	----------	----------------------------------

 - $+1.111...1 \times 2^{127} = (2 - 2^{-23}) \times 2^{127}$
 - Smallest positive number

0	00000001	00000000000000000000000000000000
---	----------	----------------------------------

 - $+1.000...0 \times 2^{-126}$
 - Negative number with smallest magnitude

1	00000001	00000000000000000000000000000000
---	----------	----------------------------------

 - $-1.000...0 \times 2^{-126}$
 - Negative number with largest magnitude

1	11111110	11111111111111111111111111111111
---	----------	----------------------------------

 - $-1.111...1 \times 2^{127} = -(2 - 2^{-23}) \times 2^{127}$

Floating-point numbers

- IEEE 754 single-precision format
 - Special numbers
 - Encoding of zero (two possible representations):

X	00000000	00000000000000000000000000000000
---	----------	----------------------------------
 - Encoding of +infinity:

0	11111111	00000000000000000000000000000000
---	----------	----------------------------------
 - Encoding of -infinity:

1	11111111	00000000000000000000000000000000
---	----------	----------------------------------
 - Encoding of NaN (result of 0/0, sqrt(-n), 0*inf, etc.):

X	11111111	Anything non-zero
---	----------	-------------------

Floating-point numbers

- IEEE 754 single-precision format
 - What about the numbers between 0 and $\pm N$ where N is the smallest representable magnitude ($1.000...0 \times 2^{-126}$)?
 - These are called denormalized numbers because they are of the form $\pm 0.b \times 2^c$ where $c = -126$
 - Largest non-negative denormalized number:

0 00000000 111111111111111111111111

– $+0.111...1 \times 2^{-126} = (1 - 2^{-23}) \times 2^{-126}$
 - Smallest positive denormalized number:

0 00000000 000000000000000000000001

– $+0.000...01 \times 2^{-126} = 2^{-149}$
 - Negative denormalized range is similar

Floating-point numbers

- IEEE 754 single-precision format
 - Representable magnitudes:
 - 0
 - Denormalized: 2^{-149} to $(1 - 2^{-23}) \times 2^{-126}$
 - Normalized: 2^{-126} to $(2 - 2^{-23}) \times 2^{127}$
 - Infinity
 - Anything with an exponent bigger than 127
 - Note: a mantissa that is larger than the largest representable mantissa (i.e., 1111...1) does not make the magnitude infinity (will discuss this soon)
 - NaN
 - Use of IEEE 754 single-precision format
 - The C data type "float" translates to 32-bit single-precision format

Floating-point numbers

- Use of IEEE 754 single-precision format
 - You can verify that the computer uses IEEE 754 single-precision format for float through the following code


```
float x = -3.75; // Can use any float value
// -3.75 = -11.11 = -1.111 x 21 in binary
// biased exponent = 128, mantissa = 111000...0

unsigned int sbit = (((unsigned int*)&x) >> 31) & 1;
unsigned int exponent = (((unsigned int*)&x) >> 23) & 0xff;

unsigned int fraction = (((unsigned int*)&x) & 0x7fffff);
printf("Sign bit: %u, biased exponent: %u, actual exponent: %d, mantissa: %#x\n", sbit, exponent, exponent - 127, fraction);
```

Floating-point numbers

- IEEE 754 double-precision format
 - Improves both precision and range over the single-precision format
 - 64-bit representation
 - MSB is sign bit
 - Least significant 52 bits represent mantissa
 - The middle 11 bits represent biased exponent with a bias of 1023
 - Actual exponent can range from -1022 to 1023
 - Largest representable normalized magnitude
 - $1.111...1 \times 2^{1023} = (2 - 2^{-52}) \times 2^{1023}$
 - Smallest representable normalized magnitude
 - $1.000...0 \times 2^{-1022} = 2^{-1022}$

Floating-point numbers

- IEEE 754 double-precision format
 - Largest representable denormalized magnitude
 - $0.111...1 \times 2^{-1022} = (1 - 2^{-52}) \times 2^{-1022}$
 - Smallest representable denormalized magnitude
 - $0.000...01 \times 2^{-1022} = 2^{-1074}$
 - Zero
 - Exponent = 0, Mantissa = 0
 - Infinity
 - Exponent = 1111111111, Mantissa = 0
 - NaN
 - Exponent = 1111111111, Mantissa = non-zero
- Use of IEEE 754 double-precision format
 - The C data type "double" translates to 64-bit double-precision format

Floating-point numbers

- Use of IEEE 754 double-precision format
 - You can verify that the computer uses IEEE 754 double-precision format for double through the following code


```
double x = -3.75; // Can use any float value
unsigned int sbit = (((unsigned long long*)&x)
>> 63) & 1;
unsigned int exponent = (((unsigned long
long*)&x) >> 52) & 0x7ff;
unsigned long long fraction = (((unsigned long
long*)&x) & 0xffffffffffff);
printf("Sign bit: %u, biased exponent: %u, actual
exponent: %d, mantissa: %#llx\n", sbit, exponent,
exponent - 1023, fraction);
```

Floating-point numbers

- IEEE 754 half-precision format
 - Lower range and precision compared to single-precision for computers having narrow busses/operands
 - 16-bit representation
 - MSB is sign bit
 - Least significant ten bits represent mantissa
 - The middle five bits represent the biased exponent with a bias of 15
 - Actual exponent can range from -14 to 15
 - Largest representable normalized magnitude
 - $1.111111111 \times 2^{15} = (2 - 1/1024) \times 2^{15}$
 - Smallest representable normalized magnitude
 - $1.0000000000 \times 2^{-14} = 2^{-14}$

Floating-point numbers

- IEEE 754 half-precision format
 - Largest representable denormalized magnitude
 - $0.111111111 \times 2^{-14} = (1 - 1/1024) \times 2^{-14}$
 - Smallest representable denormalized magnitude
 - $0.000...01 \times 2^{-14} = 2^{-24}$
 - Zero
 - Exponent = 0, Mantissa = 0
 - Infinity
 - Exponent = 11111, Mantissa = 0
 - NaN
 - Exponent = 11111, Mantissa = non-zero

Floating-point numbers

- IEEE 754 quadruple-precision format
 - Much higher range and precision compared to double-precision
 - No computer supports it yet
 - 128-bit representation
 - MSB is sign bit
 - Least significant 112 bits represent mantissa
 - The middle 15 bits represent the biased exponent with a bias of 16383
 - Actual exponent can range from -16382 to 16383
 - Largest representable normalized magnitude
 - $1.111...1 \times 2^{16383} = (2 - 2^{-112}) \times 2^{16383}$
 - Smallest representable normalized magnitude
 - $1.000...0 \times 2^{-16382} = 2^{-16382}$

Floating-point numbers

- IEEE 754 quadruple-precision format
 - Largest representable denormalized magnitude
 - $0.111...1 \times 2^{-16382} = (1 - 2^{-112}) \times 2^{-16382}$
 - Smallest representable denormalized magnitude
 - $0.000...01 \times 2^{-16382} = 2^{-16494}$
 - Zero
 - Exponent = 0, Mantissa = 0
 - Infinity
 - Exponent = 111...1, Mantissa = 0
 - NaN
 - Exponent = 111...1, Mantissa = non-zero

Overflow in floating-point ops

- What happens if the magnitude of a floating-point number exceeds the largest representable by the corresponding format?
 - E.g., float $x > (2 - 2^{-23}) \times 2^{127}$
 - May cause an overflow if exponent is larger than maximum allowed (e.g., 127 in single-precision)
 - The number is treated as +infinity or -infinity depending on the sign of the number
 - What if the fraction cannot fit within the mantissa field, but the exponent is within range?
 - Not an overflow
 - Handled by rounding the mantissa (default is round to nearest and round to nearest even for halfway round)
 - Loss of precision

Underflow in floating-point ops

- What happens if the magnitude of a floating-point number is less than the smallest representable by the corresponding format?
 - E.g., float $x < 2^{-149}$
 - May cause an underflow if an exponent smaller than the minimum (e.g., -126 in single-precision) is required to represent the number
 - E.g., 2^{-150} in single-precision or 2^{-1075} in double-precision
 - The number is treated as zero in this case
 - Denormalized numbers are said to undergo gradual underflow as more and more leading zeros appear on the right side of the binary point ultimately becoming zero below 2^{-149}

Rounding modes

- IEEE 754 rounding modes
 - Round to nearest (default behavior)
 - Round to nearest even for halfway rounding
 - After rounding, the least significant representable mantissa bit should be even (see following single-precision examples)
 - Example: 1.1111...1 (has 24 1s in mantissa) is rounded to 2.0 (in decimal) i.e., 1.0×2^1 in binary
 - Example: 1.111...101 (has 22 1s followed by a 0 and a 1 in mantissa) is rounded to 1.111...10 (has 22 1s followed by a 0 in mantissa)
 - Round toward zero
 - Round toward +infinity
 - Round toward -infinity
- A computer can choose one of the modes