

# Fundamentals of Database Systems

Assignment: 2

Due Date: 18th August, 2017

## Instructions

This question paper contains 10 questions in 6 pages.

**Q1:** Consider the following schema for an office payroll system, where the primary keys are underlined and the foreign keys are italicized.

Person(pid, fname, lname)  
Employee(pid, *desig*, salary)

Write a relational algebra query to update the salary of all employees to the average salary for their designation.

- A.  $Employee \leftarrow \pi_{pid, desig, salary=avg(salary)}(Employee * (desig \mathcal{G}_{avg(salary)}(Employee)))$
- B.  $\pi_{pid, desig, salary=avg(salary)} \sigma_{salary \neq avg(salary)}(Employee * (desig \mathcal{G}_{avg(salary)}(Employee)))$
- C.  $Employee \leftarrow \pi_{pid, salary=avg(salary)}(Employee * (desig \mathcal{G}_{avg(salary)}(Employee)))$
- D.  $Employee \leftarrow \sigma_{salary \neq avg(salary)}(Employee * (desig \mathcal{G}_{avg(salary)}(Employee)))$

**Explanation:** Updation of a set of tuples require the tuples to be selected, some function/operation on the required attribute of the selected tuple and assignment to relation without changing the schema.

**Q2:** Consider the following schema for an office payroll system, where the primary keys are underlined and the foreign keys are italicized.

Person(pid, fname, lname)  
Employee(pid, *desig*, salary)

Write an SQL query for the following relational algebra query.

$Employee \leftarrow \pi_{pid, desig, salary=avg(salary)}(Employee * (desig \mathcal{G}_{avg(salary)}(Employee)))$

- A. UPDATE Employee  
SET salary=avg(salary)  
WHERE salary != avg  
GROUP BY desig;
- B. SELECT salary=avg(salary), desig  
FROM Employee  
GROUP BY desig;
- C. WITH t as (SELECT desig as d, avg(salary) as avg  
FROM Employee  
GROUP BY desig)  
UPDATE Employee, t  
SET salary=avg  
WHERE desig = d;

D. SELECT pid, desig, salary  
FROM Employee  
WHERE salary = avg(salary)  
GROUP BY desig;

**Explanation:** Option C is correct, as it groups designations calculating average salary for each, and then updates the table appropriately.  
A is incorrect since GROUP BY cannot be used in an UPDATE query.  
B and D are incorrect because they do not perform any update, as well as logical mistakes.

**Q3:** Consider the following schema, where the primary key is underlined.

Person(pid, first\_name, last\_name, address)

Consider the following queries.

- SQL Query: SELECT first\_name, last\_name FROM Person;
- Relational Algebra Query:  $\pi_{first\_name, last\_name}(Person)$

Can the two queries have different number of tuples in their outputs?

- A. Yes
- B. No
- C. Not enough information
- D. The two queries are not the same.

**Explanation:** The two queries are same and can have different number of tuples in their answer sets based on data, since relational algebra queries output distinct tuples or a set while SQL queries are by default multi-set in nature.

**Q4:** From the following options, select the command syntax that suggests that SQL is a Data Manipulative Language (DML).

- A. INSERT INTO *table-name*  
VALUES (*comma separated list of values*);
- B. DROP TABLE *table-name*;
- C. ALTER TABLE *table-name*  
DROP *attribute-name*;
- D. None of the above

**Explanation:** INSERT command is one of the DML commands in SQL while DROP and ALTER commands are part of DDL commands.

**Q5:** From the following options, select the command syntax that suggests that SQL is a Data Definition Language (DDL).

- A. UPDATE *table-name*  
SET *attribute-name* = *value*;

- B. DELETE FROM *table-name*  
WHERE *predicate*;
- C. CREATE TABLE *table-name* (*attribute-name datatype*);
- D. All of the above

**Explanation:** CREATE TABLE is a DDL command while UPDATE and DELETE are DML commands.

**Q6:** Given below is the database schema of a hotel, where the primary keys are underlined and the foreign keys are italicized.

Room(room\_no, *intercom\_no*, tariff)  
Customer(cid, name, contact, address, staying)  
Checkin(chid, *cid*, *room\_no*, checkinTimestamp, checkoutTimestamp)

The *staying* attribute of the Customer table stores indicates if the customer is currently staying; if yes, the value is 'Yes', otherwise, it is 'No'. The system needs a trigger that changes the value of *staying* to 'No' when the check-out time is updated in the *Checkin* table. Assuming that the check-out time is updated one at a time, select the best possible answer that defines the above trigger.

- A. CREATE TRIGGER update\_staying  
AFTER UPDATE of checkoutTimestamp on Checkin  
BEGIN  
UPDATE Customer as C  
SET staying = 'No'  
WHERE C.cid = cid  
END
- B. CREATE TRIGGER update\_staying  
BEFORE INSERT on Checkin  
BEGIN  
UPDATE Customer as C  
SET staying = 'No'  
WHERE C.cid = cid  
END
- C. CREATE TRIGGER update\_staying  
BEFORE UPDATE of checkoutTimestamp on Checkin  
BEGIN  
UPDATE Customer as C  
SET staying = 'No'  
WHERE C.cid = cid  
END
- D. CREATE TRIGGER update\_staying  
AFTER INSERT on Checkin  
BEGIN  
UPDATE Customer as C  
SET staying = 'No'  
WHERE C.cid = cid  
END

**Explanation:** The query to record check-out time will be an update query and the staying status of a customer should be changed only after the check-out time has been updated from null to some value.

**Q7:** Given below is the database schema of a hotel, where the primary keys are underlined and the foreign keys are italicized.

Room(room\_no, *intercom\_no*, tariff)  
Customer(cid, name, contact, address, staying)  
Checkin(chid, *cid*, *room\_no*, checkinTimestamp, checkoutTimestamp)

What is the result of the following query?

```
UPDATE Room
SET tariff = 1.15*tariff
WHERE 1;
```

- A. Tariff of some rooms is increased by 15% for 'room\_no' 1.
- B. Tariff of 1 room is increased by 15%.
- C. Tariff of all rooms is increased by 15%.**
- D. None of the above.

**Explanation:** C is correct.

Condition '1' evaluates to TRUE, thus tariff is increased for all rooms by 15%.

**Q8:** Which is the keyword used in SQL to remove a table?

- A. REMOVE
- B. DELETE
- C. ALTER
- D. DROP**

**Explanation:** SQL specifications.

**Q9:** Consider the following database schema for an online game, where the primary keys are underlined and the foreign keys are italicized.

Account(id, *username*, email, password)  
Player(*username*, *cid*, attack-value, attack-rank, defense-value, defense-rank, gold, rank)  
Attack(aid, *attacker*, *defender*, result, gold\_stolen)  
Clan(cid, cname)

Each account is tied to a player name. Players attack each other to steal gold and raise their *attack-value*, *defense-value* to earn better rank.

Players can spend their gold to change their *attack-value*, *defense-value* at any point. Given a state of the game at any point, *attack-value* and *defense-value* can be ranked individually. Highest *attack-value* results in the *attack-rank* 1 and so on. At every hour mark, the ranks are updated based on values at that point.

Which SQL query would be used to update the *attack-rank* and *defense-rank* into the Player table?

**Note:** You may assume that "Rank()" is a function that simply gives the rank of each row of a query with respect to the other rows. Syntax of Rank() is

Rank() OVER ( [ query\_partition\_clause ] ORDER BY clause )

- A. WITH t1 as (  
     (SELECT username, attack-value AS arank FROM Player ORDER BY attack-value )  
     NATURAL JOIN  
     (SELECT username, defense-value AS drank FROM Player ORDER BY defense-value )  
   )  
   UPDATE Player  
   SET attack-rank = arank, defense-rank = drank;
- B. WITH t1 as  
     (SELECT username, attack-value AS arank FROM Player ORDER BY attack-value ),  
   t2 as  
     (SELECT username, defense-value AS drank FROM Player ORDER BY defense-value )  
   UPDATE Player  
   SET attack-rank = t1.arank, defense-rank = t2.drank;
- C. WITH t1 as (  
     (SELECT username, Rank() OVER (ORDER BY attack-value DESC) AS arank FROM Player)  
     NATURAL JOIN  
     (SELECT username, Rank() OVER (ORDER BY defense-value DESC) as drank FROM Player)  
   )  
   UPDATE Player NATURAL JOIN t1  
   SET attack-rank = arank, defense-rank = drank;
- D. WITH t1 as  
     (SELECT username, Rank() OVER (ORDER BY attack-value DESC) AS arank FROM Player)  
   t2 as  
     (SELECT username, Rank() OVER (ORDER BY defense-value DESC) as drank FROM Player)  
   UPDATE t1 NATURAL JOIN t2  
   SET attack-rank = t1.arank, defense-rank = t2.drank;

**Explanation: C**

t1 is a natural join on ordering on attack and defense values. Rank on such ordering is taken from this table and inserted in the Player table.

**Q10:** Consider the following database schema for an online game, where the primary keys are underlined and the foreign keys are italicized.

**Account**(id, *username*, email, password)

**Player**(*username*, *cid*, attack-value, attack-rank, defense-value, defense-rank, gold, rank)

**Attack**(*aid*, *attacker*, *defender*, result, gold\_stolen)

**Clan**(cid, cname)

Each account is tied to a player name. Players attack each other to steal gold and raise their *attack-value*, *defense-value* to earn better rank.

The summation of *attack-rank* and *defense-rank* decides a player's total rank, with the lowest sum yielding best rank. For example, someone having rank #1 on *attack-value* and rank #100 on *defense-value* (sum 101) will be ranked worse than someone who has rank #25 of both (sum 50).

What would be the SQL query to update current rank of the players (after the *attack-rank* and *defense-rank* has been updated)?

**Note:** You may assume that “Rank()” is a function that simply gives the rank of each row of a query with respect to the other rows. Syntax of Rank() is

Rank() OVER ( [ query\_partition\_clause] ORDER BY clause )

- A. UPDATE Player  
NATURAL JOIN  
( SELECT username, Rank() OVER (ORDER BY (attack-rank + defense-rank) ASC) AS overall\_rank FROM Player ) SET rank = overall\_rank;
- B. UPDATE Player  
SET rank = attack-rank + defense-rank;
- C. UPDATE Player  
NATURAL JOIN  
( SELECT username, Rank() OVER (ORDER BY attack-value ASC) AS a\_rank, Rank() OVER (ORDER BY defense-value ASC) AS d\_rank FROM Player )  
SET rank = a\_rank + d\_rank;
- D. UPDATE Player  
NATURAL JOIN  
( SELECT username, Rank() OVER (ORDER BY (attack-value + defense-value) ASC) AS overall\_rank FROM Player ) SET rank = overall\_rank;

**Explanation:** A is correct.

Since *attack-rank* and *defense-rank* are already present, we order on their sum and assign rank thereafter. Ordering on the *attack-value* and *defense-value* as in option D, will not be reliable since their ranks need not be directly proportional to values. B and C are clearly wrong since they are basically setting the final rank as sum of individual ranks, instead of ordering on that sum.