

CS315: DATABASE SYSTEMS CONCURRENCY CONTROL

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2nd semester, 2019-20

Tue, Wed 12:00-13:15

Locks

- Locks are used to control access to a data item
- Lock requests are made to **concurrency control manager**
- Concurrency control manager decides whether and when to grant locks
- Locking and unlocking must be atomic

- Locks are used to control access to a data item
- Lock requests are made to **concurrency control manager**
- Concurrency control manager decides whether and when to grant locks
- Locking and unlocking must be **atomic**
- A data item can be locked in two modes
 - 1 **Exclusive (X)** mode: Data item can be both written and read
 - 2 **Shared (S)** mode: Data item can only be read

- Locks are used to control access to a data item
- Lock requests are made to **concurrency control manager**
- Concurrency control manager decides whether and when to grant locks
- Locking and unlocking must be **atomic**
- A data item can be locked in two modes
 - ① **Exclusive (X)** mode: Data item can be both written and read
 - ② **Shared (S)** mode: Data item can only be read
- A lock can be granted based on the compatibility matrix
- **Lock compatibility matrix** or **conflict matrix**

	S	X
S	yes	no
X	no	no

- If a lock cannot be granted, it must wait

Locking Protocol

- A schedule must specify all locking and unlocking operations and their modes
 - $lx(a)$ requests an exclusive lock on data item a ; $ux(a)$ releases it
 - $ls(a)$ requests a shared lock on data item a ; $us(a)$ releases it
- Example: $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$

Locking Protocol

- A schedule must specify all locking and unlocking operations and their modes
 - $lx(a)$ requests an exclusive lock on data item a ; $ux(a)$ releases it
 - $ls(a)$ requests a shared lock on data item a ; $us(a)$ releases it
- Example: $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$
- Consider $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); \underline{lx_2(a)}; \underline{lx_1(b)}$

Locking Protocol

- A schedule must specify all locking and unlocking operations and their modes
 - $lx(a)$ requests an exclusive lock on data item a ; $ux(a)$ releases it
 - $ls(a)$ requests a shared lock on data item a ; $us(a)$ releases it
- Example: $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$
- Consider $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); lx_2(a); lx_1(b)$
 - **Deadlock**

Locking Protocol

- A schedule must specify all locking and unlocking operations and their modes
 - $lx(a)$ requests an exclusive lock on data item a ; $ux(a)$ releases it
 - $ls(a)$ requests a shared lock on data item a ; $us(a)$ releases it
- Example: $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$
- Consider $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); lx_2(a); lx_1(b)$
 - **Deadlock**
- **Starvation** may also happen

Locking Protocol

- A schedule must specify all locking and unlocking operations and their modes
 - $lx(a)$ requests an exclusive lock on data item a ; $ux(a)$ releases it
 - $ls(a)$ requests a shared lock on data item a ; $us(a)$ releases it
- Example: $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$
- Consider $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); lx_2(a); lx_1(b)$
 - **Deadlock**
- **Starvation** may also happen
- A **locking protocol** specifies the rules of how a transaction can acquire and release locks

Two-Phase Locking Protocol (2PL)

- Two phases
- Phase 1: **Growing (locking) phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking (unlocking) phase**
 - Transaction may release locks
 - Transaction may not obtain locks

Two-Phase Locking Protocol (2PL)

- Two phases
- Phase 1: **Growing (locking) phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking (unlocking) phase**
 - Transaction may release locks
 - Transaction may not obtain locks
- 2PL schedules are conflict serializable
 - Serialized in the order of **lock points**
 - Lock point is the time when *all* locks are obtained


Two-Phase Locking Protocol (2PL)

- Two phases
- Phase 1: **Growing (locking) phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking (unlocking) phase**
 - Transaction may release locks
 - Transaction may not obtain locks
- 2PL schedules are conflict serializable
 - Serialized in the order of **lock points**
 - Lock point is the time when *all* locks are obtained
- May suffer from deadlock
 - $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ls_2(a); lx_1(b)$

Two-Phase Locking Protocol (2PL)

- Two phases
- Phase 1: **Growing (locking) phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking (unlocking) phase**
 - Transaction may release locks
 - Transaction may not obtain locks
- 2PL schedules are conflict serializable
- Serialized in the order of **lock points**
- Lock point is the time when all locks are obtained
- May suffer from deadlock
- May suffer from cascading rollbacks

But not all conflict-serializable
schedules can be produced
by 2PL



$lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ls_2(a); lx_1(b)$

Variants of 2PL

- Basic 2PL
 - Basic protocol

Variants of 2PL

- **Basic 2PL**
 - Basic protocol
- **Strict 2PL**
 - A transaction must hold all its *exclusive* locks till it commits or aborts
 - Avoids cascading rollbacks
 - Produces strict schedules
 - May deadlock

Variants of 2PL

- Basic 2PL

- Basic protocol

- Strict 2PL

- A transaction must hold all its exclusive locks till it commits or aborts
- Avoids cascading rollbacks
- Produces strict schedules [strict recovery?]
- May deadlock

- Rigorous 2PL

- A transaction must hold *all* its locks till it commits or aborts
- Transactions can be serialized in the order of their commits

Variants of 2PL

- **Basic 2PL**

- Basic protocol

- **Strict 2PL**

- A transaction must hold all its ^Xexclusive locks till it commits or aborts
 - Avoids cascading rollbacks
 - Produces strict schedules
 - May deadlock
- No cascading rollback
- Deadlocks occur

- **Rigorous 2PL**

- A transaction must hold ^(S,X)all its locks till it commits or aborts
 - Transactions can be serialized in the order of their commits
- No cascading rollbacks

- **Conservative (static) 2PL**

- All locks are acquired atomically before a transaction begins
 - Each transaction declares its read set and write set
 - Deadlock-free
- Deadlocks occur

- No deadlock

Timestamps

- Each transaction is assigned a **timestamp** when it starts
 - Transaction T_i starting earlier has a lower timestamp than T_j starting later
- For each data item x , two timestamps are maintained
- **write-timestamp(x)**, $wts(x)$, is the largest timestamp of any transaction that executed write successfully
- **read-timestamp(x)**, $rts(x)$, is the largest timestamp of any transaction that executed read successfully
- Protocols using timestamps *cannot* deadlock

Timestamp Ordering (TO) Protocol


- Ensures that conflicting operations are executed in timestamp order
- When a transaction T requests read(x)

Timestamp Ordering (TO) Protocol

- Ensures that conflicting operations are executed in timestamp order
- When a transaction T requests $\text{read}(x)$
 - If $\text{ts}(T) > \text{wts}(x)$, no conflict

Timestamp Ordering (TO) Protocol

- Ensures that conflicting operations are executed in timestamp order
- When a transaction T requests $\text{read}(x)$
 - If $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{read}(x)$ request is executed
 - $\text{rts}(x)$ is updated to $\max\{\text{rts}(x), \text{ts}(T)\}$



Timestamp Ordering (TO) Protocol

- Ensures that conflicting operations are executed in timestamp order
- When a transaction T requests $\text{read}(x)$
 - If $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{read}(x)$ request is executed
 - $\text{rts}(x)$ is updated to $\max\{\text{rts}(x), \text{ts}(T)\}$
 - If $\text{ts}(T) < \text{wts}(x)$, then T is attempting to read a value that has been overwritten by a later transaction

Timestamp Ordering (TO) Protocol

- Ensures that conflicting operations are executed in timestamp order
- When a transaction T requests $\text{read}(x)$
 - If $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{read}(x)$ request is executed
 - $\text{rts}(x)$ is updated to $\max\{\text{rts}(x), \text{ts}(T)\}$
 - If $\text{ts}(T) < \text{wts}(x)$, then T is attempting to read a value that has been overwritten by a later transaction
 - $\text{read}(x)$ request is rejected
 - T or transaction that produced $\text{wts}(x)$ is rolled back

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{write}(x)$ request is executed
 - $\text{wts}(x)$ is updated to $\text{ts}(T)$

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{write}(x)$ request is executed
 - $\text{wts}(x)$ is updated to $\text{ts}(T)$
 - If $\text{ts}(T) < \text{rts}(x)$, then value of x that is being written should have been read earlier

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{write}(x)$ request is executed
 - $\text{wts}(x)$ is updated to $\text{ts}(T)$
 - If $\text{ts}(T) < \text{rts}(x)$, then value of x that is being written should have been read earlier
 - $\text{write}(x)$ request is rejected
 - T or transaction that produced $\text{rts}(x)$ is rolled back

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{write}(x)$ request is executed
 - $\text{wts}(x)$ is updated to $\text{ts}(T)$
 - If $\text{ts}(T) < \text{rts}(x)$, then value of x that is being written should have been read earlier
 - $\text{write}(x)$ request is rejected
 - T or transaction that produced $\text{rts}(x)$ is rolled back
 - If $\text{ts}(T) < \text{wts}(x)$, then T is attempting to write an obsolete value that has been overwritten by a later transaction

Timestamp Ordering (TO) Protocol (contd.)

- When a transaction T requests $\text{write}(x)$
 - If $\text{ts}(T) > \text{rts}(x)$ and $\text{ts}(T) > \text{wts}(x)$, no conflict
 - $\text{write}(x)$ request is executed
 - $\text{wts}(x)$ is updated to $\text{ts}(T)$
 - If $\text{ts}(T) < \text{rts}(x)$, then value of x that is being written should have been read earlier
 - $\text{write}(x)$ request is rejected
 - T or transaction that produced $\text{rts}(x)$ is rolled back
 - If $\text{ts}(T) < \text{wts}(x)$, then T is attempting to write an obsolete value that has been overwritten by a later transaction
 - $\text{write}(x)$ request is rejected
 - T or transaction producing $\text{wts}(x)$ is rolled back

Correctness

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
 - If an operation appears out of order, it is rejected

Correctness

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
 - If an operation appears out of order, it is rejected
- ✓ • No deadlock since all edges in the precedence graph are from transactions with smaller timestamp to those with larger timestamp

Correctness

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
 - If an operation appears out of order, it is rejected
- No deadlock since all edges in the precedence graph are from transactions with smaller timestamp to those with larger timestamp
- May cause starvation

Correctness

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
 - If an operation appears out of order, it is rejected
- No deadlock since all edges in the precedence graph are from transactions with smaller timestamp to those with larger timestamp
- May cause starvation
- Is not cascadeless

Correctness

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
 - If an operation appears out of order, it is rejected
- No deadlock since all edges in the precedence graph are from transactions with smaller timestamp to those with larger timestamp
- May cause starvation
- Is not cascadeless
- Is not recoverable

Modifications

- To make it recoverable
 - Use **commit dependency**
 - If T_i reads from T_j and T_j has not committed, then T_i has a commit dependency on T_j
 - Ensure that T_i does not commit before T_j commits

Modifications

- To make it recoverable
 - Use **commit dependency**
 - If T_i reads from T_j and T_j has not committed, then T_i has a commit dependency on T_j
 - Ensure that T_i does not commit before T_j commits
- To make it recoverable and cascadeless
 - All writes are performed *atomically* in the end
 - A transaction, if aborts, is re-started with a new timestamp
- To make it recoverable and cascadeless
 - Lock data that is begin written
 - Wait for it to be committed before allowing read

Modifications

- To make it recoverable
 - Use **commit dependency**
 - If T_i reads from T_j and T_j has not committed, then T_i has a commit dependency on T_j
 - Ensure that T_i does not commit before T_j commits
- To make it recoverable and cascadeless
 - All writes are performed *atomically* in the end
 - A transaction, if aborts, is re-started with a new timestamp
- To make it recoverable and cascadeless
 - Lock data that is begin written
 - Wait for it to be committed before allowing read
- **Strict timestamp ordering**: to make it strict
 - Wait for data to be committed before reading or writing

Thomas' Write Rule

- Obsolete writes may be ignored
- When T attempts to write x, if $ts(T) < wts(x)$, then T is trying to write an obsolete value of x
- **Thomas' write rule:** Rather than aborting T, ignore the write operation
 - Write is obsolete anyway

Thomas' Write Rule

- Obsolete writes may be ignored
- When T attempts to write x, if $ts(T) < wts(x)$, then T is trying to write an obsolete value of x
- **Thomas' write rule**: Rather than aborting T, ignore the write operation
 - Write is obsolete anyway
- Improves concurrency and recoverability
- Allows some view-serializable schedules that are not conflict-serializable
 - $r_1(a)w_2(a)w_1(a)w_3(a)$

Validation (Certification)-Based Protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs **validation test** to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well

Validation (Certification)-Based Protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs **validation test** to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well
- Three timestamps for each transaction
 - $start(T)$: start of execution phase
 - $validation(T)$: start of validation phase
 - $finish(T)$: end of write phase
- Timestamp of T is set to validation timestamp: $ts(T) = validation(T)$
- Serialized using this timestamp
 - Increases concurrency

Validation (Certification)-Based Protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs **validation test** to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well
- Three timestamps for each transaction
 - $start(T)$: start of execution phase
 - $validation(T)$: start of validation phase
 - $finish(T)$: end of write phase
- Timestamp of T is set to validation timestamp: $ts(T) = validation(T)$
- Serialized using this timestamp
 - Increases concurrency
- Cascadeless

Validation (Certification)-Based Protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs validation test to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well
- Three timestamps for each transaction
 - $start(T)$: start of execution phase
 - $validation(T)$: start of validation phase
 - $finish(T)$: end of write phase
- Timestamp of T is set to validation timestamp: $ts(T) = validation(T)$
- Serialized using this timestamp
 - Increases concurrency
- Cascadeless
- Starvation

Validation (Certification)-Based Protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs **validation test** to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well
- Three timestamps for each transaction
 - $start(T)$: start of execution phase
 - $validation(T)$: start of validation phase
 - $finish(T)$: end of write phase
- Timestamp of T is set to validation timestamp: $ts(T) = validation(T)$
- Serialized using this timestamp
 - Increases concurrency
- Cascadeless
- Starvation

Validation Test

- For a transaction T_i , check two conditions for all transactions T_j with $ts(T_j) < ts(T_i)$
 - $finish(T_j) < start(T_i)$
 - $finish(T_j) < validation(T_i)$ and the read-set of T_i is disjoint from the write-set of T_j
- If either of these conditions is true, validation succeeds; otherwise, it fails

Validation Test

- For a transaction T_i , check two conditions for all transactions T_j with $ts(T_j) < ts(T_i)$
 - $finish(T_j) < start(T_i)$
 - $finish(T_j) < validation(T_i)$ and the read-set of T_i is disjoint from the write-set of T_j
- If either of these conditions is true, validation succeeds; otherwise, it fails
- Justification
 - First condition ensures serial schedules
 - Writes of T_i cannot affect reads of T_j


Validation Test

- For a transaction T_i , check two conditions for all transactions T_j with $ts(T_j) < ts(T_i)$
 - $finish(T_j) < start(T_i)$
 - $finish(T_j) < validation(T_i)$ and the read-set of T_i is disjoint from the write-set of T_j
- If either of these conditions is true, validation succeeds; otherwise, it fails
- Justification
 - First condition ensures serial schedules
 - Writes of T_i cannot affect reads of T_j
 - Writes of T_j do not affect reads of T_i as they are disjoint

Multiple Granularity


- Hierarchy of data items
 - DB, Relation, Tuple, Attribute
- Locking can be done at different levels
- Locking a node *explicitly* locks all its descendants *implicitly*
 - Explicit locks
 - Implicit locks
- Granularity of locking
 - *Fine granularity*: lower in tree, high concurrency, high locking overhead
 - *Coarse granularity*: higher in tree, low concurrency, low locking overhead

Problems of Simple Locking




- Assume that T_1 has locked tuple f_2
- T_2 wants to lock a_{24}

Problems of Simple Locking




- Assume that T_1 has locked tuple f_2
- T_2 wants to lock a_{24}
 - It cannot since a_{24} is implicitly locked

Problems of Simple Locking




- Assume that T_1 has locked tuple f_2
- T_2 wants to lock a_{24}
 - It cannot since a_{24} is implicitly locked
 - Find out by traversing path from a_{24} to d
- T_3 wants to lock r_1

Problems of Simple Locking



- Assume that T_1 has locked tuple f_2
- T_2 wants to lock a_{24}
 - It cannot since a_{24} is implicitly locked
 - Find out by traversing path from a_{24} to d
- T_3 wants to lock r_1
 - It cannot since that would lock f_2 implicitly

Problems of Simple Locking



- Assume that T_1 has locked tuple f_2
- T_2 wants to lock a_{24}
 - It cannot since a_{24} is implicitly locked
 - Find out by traversing path from a_{24} to d
- T_3 wants to lock r_1
 - It cannot since that would lock f_2 implicitly
 - Find out by searching entire subtree under r_1
- Thus, for efficiency, **intention lock modes** are used
 - Ancestors of an explicitly locked node are in intention mode

Intention Lock Modes

- In addition to **shared (S)** and **exclusive (X)** locks, three additional locks
- **Intention-shared (IS)**: at least one descendant has a S lock
- **Intention-exclusive (IX)**: at least one descendant has a X lock
- **Shared and intention-exclusive (SIX)**: node is locked in S mode *and* at least one descendant has X lock

Rules of Locking

- A transaction may obtain only one lock on an entity at a time
- If two locks are needed, the more restrictive one will be acquired
- Every lock must be given notice of all lower-level locks

Rules of Locking

- A transaction may obtain only one lock on an entity at a time
- If two locks are needed, the more restrictive one will be acquired
- Every lock must be given notice of all lower-level locks
- Locks are acquired in root-to-leaf order
- Locks are released in leaf-to-root order


Rules of Locking

- A transaction may obtain only one lock on an entity at a time
- If two locks are needed, the more restrictive one will be acquired
- Every lock must be given notice of all lower-level locks
- Locks are acquired in root-to-leaf order
- Locks are released in leaf-to-root order
- Compatibility matrix

	IS	IX	S	SIX	X
IS	yes	yes	yes	yes	no
IX	yes	yes	no	no	no
S	yes	no	yes	no	no
SIX	yes	no	no	no	no
X	no	no	no	no	no

Multiple Granularity Locking Scheme

- Transaction T wants to lock a node x:
 - Lock compatibility matrix is observed
 - In S or IS mode: only if parent of x is locked by T in IX or IS mode
 - In X, SIX or IX mode: only if parent of x is locked by T in IX or SIX mode
 - Maintains 2PL, i.e., has not unlocked anything



Multiple Granularity Locking Scheme

- Transaction T wants to lock a node x :
 - Lock compatibility matrix is observed
 - In S or IS mode: only if parent of x is locked by T in IX or IS mode
 - In X, SIX or IX mode: only if parent of x is locked by T in IX or SIX mode
 - Maintains 2PL, i.e., has not unlocked anything
- Transaction T wants to unlock a node x :
 - No child of x is currently locked by T

Multiple Granularity Locking Scheme

- Transaction T wants to lock a node x :
 - Lock compatibility matrix is observed
 - In S or IS mode: only if parent of x is locked by T in IX or IS mode
 - In X, SIX or IX mode: only if parent of x is locked by T in IX or SIX mode
 - Maintains 2PL, i.e., has not unlocked anything
- Transaction T wants to unlock a node x :
 - No child of x is currently locked by T
- Ensures conflict serializability

SIX Lock

- Suppose T_1 wants to read r_1 but only modify a_{24}
- Locking r_1 in IX mode will allow other transactions to lock r_1 in IX mode
 - Unsafe as T_1 is reading r_1
- Locking r_1 in S mode will allow other transactions to lock r_1 in S mode and read everything
 - Unsafe as T_1 is modifying a_{24}
- SIX lock compromises and is safer

Example

- T_1 wants to read a_{12}

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d, r_1 in IS mode and f_1 in S mode
- T_4 wants to read d

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d, r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d, r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode
- T_1 and T_2

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d, r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode
- T_1 and T_2 can execute concurrently
- T_1, T_3 and T_4

Example

- T_1 wants to read a_{12}
 - Locks d, r_1, f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d, r_1, f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d, r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode
- T_1 and T_2 can execute concurrently
- T_1, T_3 and T_4 can execute concurrently
- T_2 and T_3

Example

- T_1 wants to read a_{12}
 - Locks d , r_1 , f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d , r_1 , f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d , r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode
- T_1 and T_2 can execute concurrently
- T_1 , T_3 and T_4 can execute concurrently
- T_2 and T_3 cannot execute concurrently
- T_2 and T_4

Example

- T_1 wants to read a_{12}
 - Locks d , r_1 , f_1 in IS mode and a_{12} in S mode
- T_2 wants to write a_{14}
 - Locks d , r_1 , f_1 in IX mode and a_{12} in X mode
- T_3 wants to read f_1
 - Locks d , r_1 in IS mode and f_1 in S mode
- T_4 wants to read d
 - Locks d in S mode
- T_1 and T_2 can execute concurrently
- T_1 , T_3 and T_4 can execute concurrently
- T_2 and T_3 cannot execute concurrently
- T_2 and T_4 cannot execute concurrently

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
 - Older transactions wait for younger ones to release data item (**wait**)
 - Younger ones do not wait, they roll back (**die**)
 - Young transactions may die many times

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
 - Older transactions wait for younger ones to release data item (**wait**)
 - Younger ones do not wait, they roll back (**die**)
 - Young transactions may die many times
- **Wound-wait**: Preemptive
 - Older transactions kill younger ones and force them to release data item (**wound**)
 - Younger ones wait (**wait**)

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
 - Older transactions wait for younger ones to release data item (**wait**)
 - Younger ones do not wait, they roll back (**die**)
 - Young transactions may die many times
- **Wound-wait**: Preemptive
 - Older transactions kill younger ones and force them to release data item (**wound**)
 - Younger ones wait (**wait**)
- Transactions are re-started with the *same* timestamp

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
 - Older transactions wait for younger ones to release data item (**wait**)
 - Younger ones do not wait, they roll back (**die**)
 - Young transactions may die many times
- **Wound-wait**: Preemptive
 - Older transactions kill younger ones and force them to release data item (**wound**)
 - Younger ones wait (**wait**)
- Transactions are re-started with the *same* timestamp
- No starvation

Deadlock Prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
 - Older transactions wait for younger ones to release data item (**wait**)
 - Younger ones do not wait, they roll back (**die**)
 - Young transactions may die many times
- **Wound-wait**: Preemptive
 - Older transactions kill younger ones and force them to release data item (**wound**)
 - Younger ones wait (**wait**)
- Transactions are re-started with the *same* timestamp
- No starvation
- Wound-wait has fewer rollbacks than wait-die
 - Less likely for old transactions to not finish and want a lock from a young transaction



Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress
 - One inducing least number of cascading rollbacks

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress
 - One inducing least number of cascading rollbacks
- How far to rollback?

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress
 - One inducing least number of cascading rollbacks
- How far to rollback?
 - **Total rollback**: Completely abort and re-start
 - **Partial rollback**: Rollback to only as far as necessary to break deadlock

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress
 - One inducing least number of cascading rollbacks
- How far to rollback?
 - **Total rollback**: Completely abort and re-start
 - **Partial rollback**: Rollback to only as far as necessary to break deadlock
- Starvation happens when the same transaction is repeatedly chosen as the victim

Deadlock Recovery

- Deadlocks can be detected by a **wait-for** graph
 - T_i waits for T_j , T_j waits for T_k , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
 - One with lowest cost
 - One with least progress
 - One inducing least number of cascading rollbacks
- How far to rollback?
 - **Total rollback**: Completely abort and re-start
 - **Partial rollback**: Rollback to only as far as necessary to break deadlock
- Starvation happens when the same transaction is repeatedly chosen as the victim
 - Factor number of rollbacks when choosing victim

Insert and Delete

- **insert(x)**: inserts the data item x
- **delete(x)**: deletes the data item x
- Logical errors
 - read(x), write(x) before insert(x)
 - read(x), write(x) after delete(x)
 - delete(x) after delete(x)
 - insert(x) after insert(x)
- Conflicts
 - Similar to write(x)
 - Conflicts with operations on relation

delete(x) before insert(x)

Phantom Phenomenon

- Transaction T1 reads an entire relation to compute an aggregate
- Transaction T2 inserts tuples to the relation
- Transactions T1 and T2 *logically conflict*

Phantom Phenomenon

- Transaction T1 reads an entire relation to compute an aggregate
- Transaction T2 inserts tuples to the relation
- Transactions T1 and T2 *logically conflict*
- However, they do not conflict on any tuple
- They conflict on a **phantom tuple**
 - This is called **phantom phenomenon**

Phantom Phenomenon

- Transaction T1 reads an entire relation to compute an aggregate
- Transaction T2 inserts tuples to the relation
- Transactions T1 and T2 *logically conflict*
- However, they do not conflict on any tuple
- They conflict on a **phantom tuple**
 - This is called **phantom phenomenon**
- Exclusive lock on relation solves this
- Multi-granularity locking

Phantom Phenomenon

- Transaction T1 reads an entire relation to compute an aggregate
- Transaction T2 inserts tuples to the relation
- Transactions T1 and T2 *logically conflict*
- However, they do not conflict on any tuple
- They conflict on a **phantom tuple**
 - This is called **phantom phenomenon**
- Exclusive lock on relation solves this
- Multi-granularity locking
- If index structure is used, **index locking** protocol improves concurrency by locking index nodes
 - Avoids phantom phenomenon since every transaction needs to lock all accessed nodes