

CS315: DATABASE SYSTEMS

RELATIONAL MODEL

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2nd semester, 2019-20

Tue, Wed 12:00-13:15

Relation

- A **relation** is a subset of the cross-product of sets
- For sets D_1, D_2, \dots, D_n , a relation r is a set of n -ary tuples of the form (a_1, a_2, \dots, a_n) where each $a_i \in D_i$
- Example
 - `name = {A, B, C}`
 - `designation = {L, E, W}`
 - `identifier = \mathcal{N}`
 - $r = \{(A, E, 4), (B, E, 9), (C, W, 23)\}$ is a relation over
`name × designation × identifier`
- Relations are *unordered*

Relation

- A **relation** is a subset of the cross-product of sets
- For sets D_1, D_2, \dots, D_n , a relation r is a set of n -ary tuples of the form (a_1, a_2, \dots, a_n) where each $a_i \in D_i$
- Example
 - `name` = {A, B, C}
 - `designation` = {L, E, W}
 - `identifier` = \mathcal{N}
 - $r = \{(A, E, 4), (B, E, 9), (C, W, 23)\}$ is a relation over
`name` \times `designation` \times `identifier`
- Relations are *unordered*
- Generally depicted as a table

name	designation	identifier
A	E	4
B	E	9
C	W	23

Attribute

- Each attribute of a relation has a name
- There is a domain for each attribute
- Attributes are *generally* **atomic**
 - Indivisible, not sets
- Domain is atomic if all members are atomic
- Special value **null** in every domain

Relation Schema and Tuple

- The sets define a **relation schema**
- Example
 - Schema is `Person = (name, designation, identifier)`
- Relations are defined over a schema
- If schema is R , relation is denoted by $r(R)$
 - Example: `persons(Person)`
- A **relation instance** is a particular instance from the schema
 - Earlier example
- An element of a relation (instance) is a **tuple**

Handwritten diagram illustrating a relation instance. A table is shown with three rows and three columns. The first column contains identifiers (1, 2, 3), the second column contains names (A, B, C), and the third column contains designations (1st, 2nd, 3rd). A yellow arrow points from the word "Tuple" to the second row of the table, which contains the values (2, B, 2nd).

1.	A	1st
2.	B	2nd
3	C	3rd

Relation Schema and Tuple

- The sets define a **relation schema**
- Example
 - Schema is `Person = (name, designation, identifier)`
- Relations are defined over a schema
- If schema is R , relation is denoted by $r(R)$
 - Example: `persons(Person)`
- A **relation instance** is a particular instance from the schema
 - Earlier example
- An element of a relation (instance) is a **tuple**
- Tuples are rows and attributes are columns

Database

- Consists of multiple *inter-related* relations
- Each relation stores information about a particular relationship

Database

- Consists of multiple *inter-related* relations
- Each relation stores information about a particular relationship
- Alternatively, a single relation can store all data
- Problems
 - Data repetition
 - Need for null values
- Normalization theory deals with how to design relation schemas

Key

- $K \subseteq R$ is a **superkey** of R if and only if values for K are sufficient to identify a unique tuple in *all* possible relations $r(R)$
 - Possible $r(R)$ signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
 - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey

Key

- $K \subseteq R$ is a **superkey** of R if and only if values for K are sufficient to identify a unique tuple in *all possible relations* $r(R)$
 - Possible $r(R)$ signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
 - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey K is a **candidate key** if K is **minimal**, i.e., no proper subset of it is a superkey
 - {name} is a candidate key

Key

- $K \subseteq R$ is a **superkey** of R if and only if values for K are sufficient to identify a unique tuple in *all* possible relations $r(R)$
 - Possible $r(R)$ signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
 - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey K is a **candidate key** if K is minimal, i.e., no proper subset of it is a superkey
 - {name} is a candidate key
- There may be multiple candidate keys
 - {name}, {identifier} are candidate keys

Key

- $K \subseteq R$ is a **superkey** of R if and only if values for K are sufficient to identify a unique tuple in *all* possible relations $r(R)$
 - Possible $r(R)$ signifies a relation that can exist from the data that is being modeled
- Example: {name} is a superkey if each person has a unique name, otherwise not
- All supersets of superkeys are superkeys
 - {name, designation} is also a superkey
- In practical situations, an id will be used which is guaranteed to be a superkey
- A superkey K is a **candidate key** if K is minimal, i.e., no proper subset of it is a superkey
 - {name} is a candidate key
- There may be multiple candidate keys
 - {name}, {identifier} are candidate keys
- **Primary key** is a candidate key chosen to serve as the primary means of identifying tuples
 - Choice is arbitrary as it depends on the database designer
 - Other candidate keys are called **secondary keys**

Foreign Key

- A relation schema may have an attribute that is unique (e.g., a primary key) in another schema
- This attribute is then called a **foreign key**
- Example
 - depositor = (name, number)
 - customer = (name, street, city)
 - account = (number, balance)
 - name and number in depositor are foreign keys
- Values in the foreign key attribute of the **referencing relation** may only come from those in the primary key of the **referenced relation**

Entity-Relationship Model

- **Entity-Relationship Model:** A database is a collection of *entities* and *relationships* among them

Entity-Relationship Model

- **Entity-Relationship Model**: A database is a collection of *entities* and *relationships* among them
- An **entity** is an object that exists and is distinguishable from objects, e.g., person “A”
- Entities have **attributes**, e.g., persons have names
- An **entity set** is a set of entities of the same type that share the same properties, e.g., a set of persons

Entity-Relationship Model

- **Entity-Relationship Model**: A database is a collection of *entities* and *relationships* among them
- An **entity** is an object that exists and is distinguishable from objects, e.g., person “A”
- Entities have **attributes**, e.g., persons have names
- An **entity set** is a set of entities of the same type that share the same properties, e.g., a set of persons
- A **relationship** is an association among the entities
- A **relationship set** is a relation among two or more entities, each taken from an entity set, e.g., person is a depositor for an account

Entity-Relationship Model

- **Entity-Relationship Model**: A database is a collection of *entities* and *relationships* among them
- An **entity** is an object that exists and is distinguishable from objects, e.g., person “A”
- Entities have **attributes**, e.g., persons have names
- An **entity set** is a set of entities of the same type that share the same properties, e.g., a set of persons
- A **relationship** is an association among the entities
- A **relationship set** is a relation among two or more entities, each taken from an entity set, e.g., person is a depositor for an account
- Relationship sets can also have attributes, e.g., access date for depositor
- Primary keys of entity sets form a superkey of the relationship set

Model Parameters

- Relationships between two entity sets are *binary*
 - It is rare to have more than degree two
- Cardinality constraints for binary relationships can be one-to-one, one-to-many, many-to-one or many-to-many

Model Parameters

- Relationships between two entity sets are *binary*
 - It is rare to have more than degree two
- Cardinality constraints for binary relationships can be one-to-one, one-to-many, many-to-one or many-to-many
- Attributes may be
 - Simple or composite
 - Single-valued or multi-valued (e.g., phone number)
 - Derived (e.g., age)

Model Parameters

- Relationships between two entity sets are *binary*
 - It is rare to have more than degree two
- Cardinality constraints for binary relationships can be one-to-one, one-to-many, many-to-one or many-to-many
- Attributes may be
 - Simple or composite
 - Single-valued or multi-valued (e.g., phone number)
 - Derived (e.g., age)
- Participation can be total (every entity participates in the relationship) or partial (not all participate)

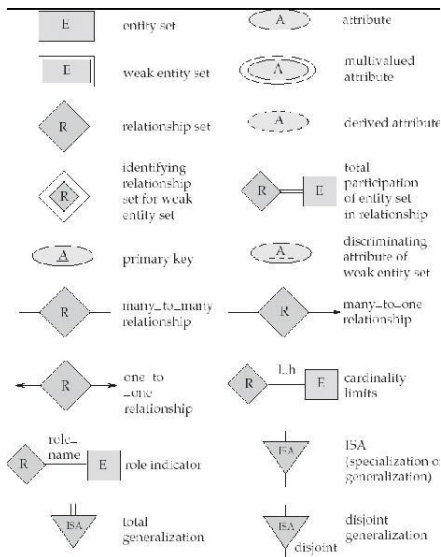
Model Parameters

- Relationships between two entity sets are *binary*
 - It is rare to have more than degree two
- Cardinality constraints for binary relationships can be one-to-one, one-to-many, many-to-one or many-to-many
- Attributes may be
 - Simple or composite
 - Single-valued or multi-valued (e.g., phone number)
 - Derived (e.g., age)
- Participation can be total (every entity participates in the relationship) or partial (not all participate)
- Specialization, generalization, aggregation are features of extended ER model

Weak Entity Sets

- An entity set that does not have a primary key is called a weak entity set
- Its existence depends on the existence of another entity set called the identifying entity set or owner entity set
- The identifying relationship set that exists between the two must be total and many-to-one from the weak entity set
- A weak entity set has a discriminator or partial key instead of a primary key
- The discriminator distinguishes weak entities that are related to the same entity of the identifying entity set
- The primary key is formed by the primary key of the identifying set and the discriminator

ER Diagram: Summary

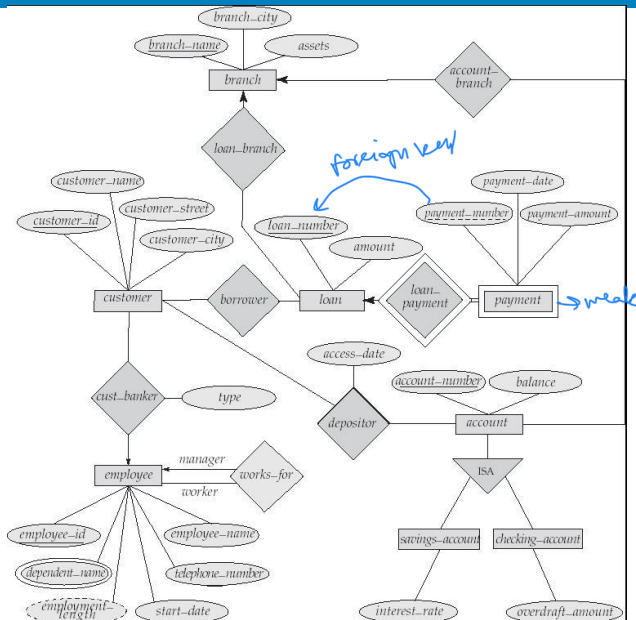


ER Diagram: Description

- Entity sets: rectangles
- Relationship sets: diamonds
- Attributes: ellipses
 - Multivalued attributes: double ellipses
 - Derived attributes: dashed ellipses
- Primary keys: underlines
- Roles: on links
- Cardinality constraints
 - One: directed
 - Many: undirected
 - One-to-many: directed-diamond-undirected
- Participation
 - Total: double line
 - Partial: single line
- Cardinality limits: on lines
- Weak entity sets: double rectangle
- Weak relationship set: double ellipse
- Discriminator of a weak entity sets: underline with dashed lines



Example: Banking Schema



Reduction to Relational Model

ER Model

- Entity sets and relationships sets are reduced uniformly to relations → Relational Model
- A weak entity set is reduced to a relation by including the primary key attributes of the identifying set
 - A foreign key relationship is established
- Many-to-one and one-to-many relationships that are total on the many side may not be reduced to a relation
 - Primary key of entity on “one” side is added to relation of entity on “many” side
- For one-to-one relationships,

Reduction to Relational Model

- Entity sets and relationships sets are reduced uniformly to relations
- A weak entity set is reduced to a relation by including the primary key attributes of the identifying set
 - A foreign key relationship is established
- Many-to-one and one-to-many relationships that are total on the many side may not be reduced to a relation
 - Primary key of entity on “one” side is added to relation of entity on “many” side
- For one-to-one relationships, any side can be chosen as “many”
- If participation is partial,

Reduction to Relational Model

- Entity sets and relationships sets are reduced uniformly to relations
- A weak entity set is reduced to a relation by including the primary key attributes of the identifying set
 - A foreign key relationship is established
- Many-to-one and one-to-many relationships that are total on the many side may not be reduced to a relation
 - Primary key of entity on “one” side is added to relation of entity on “many” side
- For one-to-one relationships, any side can be chosen as “many”
- If participation is partial, null values are used
- Many-to-many relationships

Reduction to Relational Model

- Entity sets and relationships sets are reduced uniformly to relations
- A weak entity set is reduced to a relation by including the primary key attributes of the identifying set
 - A foreign key relationship is established
- Many-to-one and one-to-many relationships that are total on the many side may not be reduced to a relation
 - Primary key of entity on “one” side is added to relation of entity on “many” side
- For one-to-one relationships, any side can be chosen as “many”
- If participation is partial, null values are used
- Many-to-many relationships must be reduced to relations
- Each component of a composite attribute is modeled separately
- Multivalued attributes

Reduction to Relational Model

- Entity sets and relationships sets are reduced uniformly to relations
- A weak entity set is reduced to a relation by including the primary key attributes of the identifying set
 - A foreign key relationship is established
- Many-to-one and one-to-many relationships that are total on the many side may not be reduced to a relation
 - Primary key of entity on “one” side is added to relation of entity on “many” side
- For one-to-one relationships, any side can be chosen as “many”
- If participation is partial, null values are used
- Many-to-many relationships must be reduced to relations
- Each component of a composite attribute is modeled separately
- Multivalued attributes are reduced to relations that include the primary key of the entity set