

Q1: Objdump 1

a)

main.o

main.o: file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <main>:

```
0: f3 0f 1e fa      endbr64
4: 55              push %rbp
5: 48 89 e5        mov %rsp,%rbp
8: b8 00 00 00 00   mov $0x0,%eax
d: e8 00 00 00 00   callq 12 <main+0x12>
               e: R_X86_64_PLT32   swap-0x4
12: 8b 15 00 00 00 00   mov 0x0(%rip),%edx   # 18 <main+0x18>
               14: R_X86_64_PC32   buf
18: 8b 05 00 00 00 00   mov 0x0(%rip),%eax   # 1e <main+0x1e>
               1a: R_X86_64_PC32   buf-0x4
1e: 89 c6          mov %eax,%esi
20: 48 8d 3d 00 00 00 00   lea 0x0(%rip),%rdi   # 27 <main+0x27>
               23: R_X86_64_PC32   .rodata-0x4
27: b8 00 00 00 00   mov $0x0,%eax
2c: e8 00 00 00 00   callq 31 <main+0x31>
               2d: R_X86_64_PLT32   printf-0x4
31: b8 00 00 00 00   mov $0x0,%eax
36: 5d              pop %rbp
37: c3              retq
```

Disassembly of section .data:

0000000000000000 <buf>:

```
0: 72 00          jb 2 <buf+0x2>
2: 00 00          add %al,(%rax)
4: 56            push %rsi
5: 00 00          add %al,(%rax)
```

...

Disassembly of section .rodata:

0000000000000000 <..rodata>:

```

0: 62          (bad)
1: 75 66      jne 69 <main+0x69>
3: 5b         pop %rbx
4: 30 5d 3d   xor %bl,0x3d(%rbp)
7: 20 25 64 20 62 75 and %ah,0x75622064(%rip) # 75622071
<main+0x75622071>
d: 66 5b     pop %bx
f: 31 5d 3d   xor %ebx,0x3d(%rbp)
12: 20        .byte 0x20
13: 25        .byte 0x25
14: 64 0a 00   or %fs:(%rax),%al

```

Disassembly of section .comment:

```

0000000000000000 <.comment>:
0: 00 47 43    add %al,0x43(%rdi)
3: 43 3a 20    rex.XB cmp (%r8),%spl
6: 28 55 62    sub %dl,0x62(%rbp)
9: 75 6e      jne 79 <main+0x79>
b: 74 75      je 82 <main+0x82>
d: 20 39      and %bh,(%rcx)
f: 2e 32 2e   xor %cs:(%rsi),%ch
12: 31 2d 39 75 62 75 xor %ebp,0x75627539(%rip) # 75627551
<main+0x75627551>
18: 6e         outsb %ds:(%rsi),(%dx)
19: 74 75      je 90 <main+0x90>
1b: 32 29      xor (%rcx),%ch
1d: 20 39      and %bh,(%rcx)
1f: 2e 32 2e   xor %cs:(%rsi),%ch
22: 31 20      xor %esp,(%rax)
24: 32 30      xor (%rax),%dh
26: 31 39      xor %edi,(%rcx)
28: 31 30      xor %esi,(%rax)
2a: 30 38      xor %bh,(%rax)
...

```

Disassembly of section .note.gnu.property:

```

0000000000000000 <.note.gnu.property>:
0: 04 00      add $0x0,%al
2: 00 00      add %al,(%rax)
4: 10 00      adc %al,(%rax)
6: 00 00      add %al,(%rax)
8: 05 00 00 00 47 add $0x47000000,%eax

```

```

d: 4e 55      rex.WRX push %rbp
f: 00 02      add  %al,(%rdx)
11: 00 00      add  %al,(%rax)
13: c0 04 00 00  rolb $0x0,(%rax,%rax,1)
17: 00 03      add  %al,(%rbx)
19: 00 00      add  %al,(%rax)
1b: 00 00      add  %al,(%rax)
1d: 00 00      add  %al,(%rax)
...

```

Disassembly of section .eh_frame:

```

0000000000000000 <.eh_frame>:
0: 14 00      adc  $0x0,%al
2: 00 00      add  %al,(%rax)
4: 00 00      add  %al,(%rax)
6: 00 00      add  %al,(%rax)
8: 01 7a 52   add  %edi,0x52(%rdx)
b: 00 01      add  %al,(%rcx)
d: 78 10      js   1f< .eh_frame+0x1f>
f: 01 1b      add  %ebx,(%rbx)
11: 0c 07      or   $0x7,%al
13: 08 90 01 00 00 1c  or   %dl,0x1c000001(%rax)
19: 00 00      add  %al,(%rax)
1b: 00 1c 00   add  %bl,(%rax,%rax,1)
1e: 00 00      add  %al,(%rax)
20: 00 00      add  %al,(%rax)
    20: R_X86_64_PC32    .text
22: 00 00      add  %al,(%rax)
24: 38 00      cmp  %al,(%rax)
26: 00 00      add  %al,(%rax)
28: 00 45 0e   add  %al,0xe(%rbp)
2b: 10 86 02 43 0d 06  adc  %al,0x60d4302(%rsi)
31: 6f        outsl %ds:(%rsi),(%dx)
32: 0c 07      or   $0x7,%al
34: 08 00      or   %al,(%rax)

```

swap.o

swap.o: file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <swap>:

```
0: f3 0f 1e fa      endbr64
4: 55              push  %rbp
5: 48 89 e5         mov   %rsp,%rbp
8: 48 8d 05 00 00 00 lea    0x0(%rip),%rax    # f <swap+0xf>
    b: R_X86_64_PC32    buf
f: 48 89 05 00 00 00 mov   %rax,0x0(%rip)    # 16 <swap+0x16>
    12: R_X86_64_PC32    bufp1-0x4
16: 48 8b 05 00 00 00 mov   0x0(%rip),%rax    # 1d <swap+0x1d>
    19: R_X86_64_PC32    bufp0-0x4
1d: 8b 00           mov   (%rax),%eax
1f: 89 05 00 00 00 00 mov   %eax,0x0(%rip)    # 25 <swap+0x25>
    21: R_X86_64_PC32    .data-0x4
25: 48 8b 15 00 00 00 mov   0x0(%rip),%rdx    # 2c <swap+0x2c>
    28: R_X86_64_PC32    bufp1-0x4
2c: 48 8b 05 00 00 00 mov   0x0(%rip),%rax    # 33 <swap+0x33>
    2f: R_X86_64_PC32    bufp0-0x4
33: 8b 12           mov   (%rdx),%edx
35: 89 10           mov   %edx,(%rax)
37: 48 8b 05 00 00 00 mov   0x0(%rip),%rax    # 3e <swap+0x3e>
    3a: R_X86_64_PC32    bufp1-0x4
3e: 8b 15 00 00 00 00 mov   0x0(%rip),%edx    # 44 <swap+0x44>
    40: R_X86_64_PC32    .data-0x4
44: 89 10           mov   %edx,(%rax)
46: 90              nop
47: 5d              pop   %rbp
48: c3              retq
```

Disassembly of section .data:

0000000000000000 <temp.1915>:

```
0: 64 00 00        add   %al,%fs:(%rax)
...
```

Disassembly of section .data.rel:

0000000000000000 <bufp0>:

```
...
0: R_X86_64_64    buf
```

Disassembly of section .comment:

0000000000000000 <.comment>:

```
0: 00 47 43        add   %al,0x43(%rdi)
```

```

3: 43 3a 20      rex.XB cmp (%r8),%spl
6: 28 55 62      sub  %dl,0x62(%rbp)
9: 75 6e         jne  79 <swap+0x79>
b: 74 75         je   82 <swap+0x82>
d: 20 39         and  %bh,(%rcx)
f: 2e 32 2e     xor  %cs:(%rsi),%ch
12: 31 2d 39 75 62 75  xor  %ebp,0x75627539(%rip)    # 75627551
<swap+0x75627551>
18: 6e           outsb %ds:(%rsi),(%dx)
19: 74 75         je   90 <swap+0x90>
1b: 32 29         xor  (%rcx),%ch
1d: 20 39         and  %bh,(%rcx)
1f: 2e 32 2e     xor  %cs:(%rsi),%ch
22: 31 20         xor  %esp,(%rax)
24: 32 30         xor  (%rax),%dh
26: 31 39         xor  %edi,(%rcx)
28: 31 30         xor  %esi,(%rax)
2a: 30 38         xor  %bh,(%rax)
...

```

Disassembly of section .note.gnu.property:

```

0000000000000000 <.note.gnu.property>:
0: 04 00         add  $0x0,%al
2: 00 00         add  %al,(%rax)
4: 10 00         adc  %al,(%rax)
6: 00 00         add  %al,(%rax)
8: 05 00 00 00 47  add  $0x47000000,%eax
d: 4e 55         rex.WRX push %rbp
f: 00 02         add  %al,(%rdx)
11: 00 00         add  %al,(%rax)
13: c0 04 00 00    rolb $0x0,(%rax,%rax,1)
17: 00 03         add  %al,(%rbx)
19: 00 00         add  %al,(%rax)
1b: 00 00         add  %al,(%rax)
1d: 00 00         add  %al,(%rax)
...

```

Disassembly of section .eh_frame:

```

0000000000000000 <.eh_frame>:
0: 14 00         adc  $0x0,%al
2: 00 00         add  %al,(%rax)
4: 00 00         add  %al,(%rax)

```

```

6: 00 00      add  %al,(%rax)
8: 01 7a 52    add  %edi,0x52(%rdx)
b: 00 01      add  %al,(%rcx)
d: 78 10      js   1f <.eh_frame+0x1f>
f: 01 1b      add  %ebx,(%rbx)
11: 0c 07      or   $0x7,%al
13: 08 90 01 00 00 1c  or  %dl,0x1c000001(%rax)
19: 00 00      add  %al,(%rax)
1b: 00 1c 00    add  %bl,(%rax,%rax,1)
1e: 00 00      add  %al,(%rax)
20: 00 00      add  %al,(%rax)
    20: R_X86_64_PC32    .text
22: 00 00      add  %al,(%rax)
24: 49 00 00    rex.WB add %al,(%r8)
27: 00 00      add  %al,(%rax)
29: 45 0e      rex.RB (bad)
2b: 10 86 02 43 0d 06  adc  %al,0x60d4302(%rsi)
31: 02 40 0c    add  0xc(%rax),%al
34: 07         (bad)
35: 08 00      or   %al,(%rax)

```

other.o

other.o: file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <f>:

```

0: f3 0f 1e fa    endbr64
4: 55             push  %rbp
5: 48 89 e5       mov  %rsp,%rbp
8: b8 00 00 00 00  mov  $0x0,%eax
d: e8 03 00 00 00  callq 15 <sf>
12: 90            nop
13: 5d            pop   %rbp
14: c3            retq

```

0000000000000015 <sf>:

```

15: f3 0f 1e fa    endbr64
19: 55             push  %rbp
1a: 48 89 e5       mov  %rsp,%rbp
1d: c7 05 00 00 00 00 03  movl  $0x3,0x0(%rip)    # 27 <sf+0x12>
24: 00 00 00
    1f: R_X86_64_PC32    buf-0x8

```

```

27: c7 05 00 00 00 00 04  movl  $0x4,0x0(%rip)    # 31 <sf+0x1c>
2e: 00 00 00
      29: R_X86_64_PC32      buf-0x4
31: 90                      nop
32: 5d                      pop  %rbp
33: c3                      retq

```

Disassembly of section .comment:

0000000000000000 <.comment>:

```

0: 00 47 43      add  %al,0x43(%rdi)
3: 43 3a 20      rex.XB cmp (%r8),%spl
6: 28 55 62      sub  %dl,0x62(%rbp)
9: 75 6e        jne  79 <sf+0x64>
b: 74 75        je   82 <sf+0x6d>
d: 20 39        and  %bh,(%rcx)
f: 2e 32 2e     xor  %cs:(%rsi),%ch
12: 31 2d 39 75 62 75  xor  %ebp,0x75627539(%rip)    # 75627551 <sf+0x7562753c>
18: 6e          outsb %ds:(%rsi),(%dx)
19: 74 75        je   90 <sf+0x7b>
1b: 32 29        xor  (%rcx),%ch
1d: 20 39        and  %bh,(%rcx)
1f: 2e 32 2e     xor  %cs:(%rsi),%ch
22: 31 20        xor  %esp,(%rax)
24: 32 30        xor  (%rax),%dh
26: 31 39        xor  %edi,(%rcx)
28: 31 30        xor  %esi,(%rax)
2a: 30 38        xor  %bh,(%rax)
...

```

Disassembly of section .note.gnu.property:

0000000000000000 <.note.gnu.property>:

```

0: 04 00      add  $0x0,%al
2: 00 00      add  %al,(%rax)
4: 10 00      adc  %al,(%rax)
6: 00 00      add  %al,(%rax)
8: 05 00 00 00 47  add  $0x47000000,%eax
d: 4e 55      rex.WRX push %rbp
f: 00 02      add  %al,(%rdx)
11: 00 00      add  %al,(%rax)
13: c0 04 00 00    rolb $0x0,(%rax,%rax,1)
17: 00 03      add  %al,(%rbx)
19: 00 00      add  %al,(%rax)

```

```

1b: 00 00      add  %al,(%rax)
1d: 00 00      add  %al,(%rax)
...

```

Disassembly of section .eh_frame:

```

0000000000000000 <.eh_frame>:
 0: 14 00      adc  $0x0,%al
 2: 00 00      add  %al,(%rax)
 4: 00 00      add  %al,(%rax)
 6: 00 00      add  %al,(%rax)
 8: 01 7a 52   add  %edi,0x52(%rdx)
 b: 00 01      add  %al,(%rcx)
 d: 78 10     js   1f<.eh_frame+0x1f>
 f: 01 1b     add  %ebx,(%rbx)
11: 0c 07     or   $0x7,%al
13: 08 90 01 00 00 1c or   %dl,0x1c000001(%rax)
19: 00 00      add  %al,(%rax)
1b: 00 1c 00   add  %bl,(%rax,%rax,1)
1e: 00 00      add  %al,(%rax)
20: 00 00      add  %al,(%rax)
    20: R_X86_64_PC32 .text
22: 00 00      add  %al,(%rax)
24: 15 00 00 00 00   adc  $0x0,%eax
29: 45 0e     rex.RB (bad)
2b: 10 86 02 43 0d 06 adc  %al,0x60d4302(%rsi)
31: 4c 0c 07   rex.WR or $0x7,%al
34: 08 00     or   %al,(%rax)
36: 00 00      add  %al,(%rax)
38: 1c 00     sbb  $0x0,%al
3a: 00 00      add  %al,(%rax)
3c: 3c 00     cmp  $0x0,%al
3e: 00 00      add  %al,(%rax)
40: 00 00      add  %al,(%rax)
    40: R_X86_64_PC32 .text+0x15
42: 00 00      add  %al,(%rax)
44: 1f       (bad)
45: 00 00      add  %al,(%rax)
47: 00 00      add  %al,(%rax)
49: 45 0e     rex.RB (bad)
4b: 10 86 02 43 0d 06 adc  %al,0x60d4302(%rsi)
51: 56       push %rsi
52: 0c 07     or   $0x7,%al
54: 08 00     or   %al,(%rax)

```


a) main.o :

```
.text
  1. 1e: R_X86_64_PLT32    swap-0x4
  2. 14: R_X86_64_PC32     buf
  3. 1a: R_X86_64_PC32     buf-0x4
  4. 23: R_X86_64_PC32     .rodata-0x4
  5. 2d: R_X86_64_PLT32    printf-0x4
.eh_frame
  1. 20: R_X86_64_PC32     .text+0x0
```

swap.o :

```
.text
  1. b: R_X86_64_PC32      buf
  2. 12: R_X86_64_PC32     bufp1-0x4
  3. 19: R_X86_64_PC32     bufp0-0x4
  4. 21: R_X86_64_PC32     .data-0x4
  5. 28: R_X86_64_PC32     bufp1-0x4
  6. 2f: R_X86_64_PC32     bufp0-0x4
  7. 3a: R_X86_64_PC32     bufp1-0x4
  8. 40: R_X86_64_PC32     .data-0x4
.data.rel
  1. 0: R_X86_64_64      buf
.eh_frame
  1. 20: R_X86_64_PC32     .text+0x0
```

other.o :

```
.text
  1. 1f: R_X86_64_PC32     buf-0x8
  2. 29: R_X86_64_PC32     buf-0x4
.eh_frame
  1. 20: R_X86_64_PC32     .text+0x0
  2. 40: R_X86_64_PC32     .text+0x15
```

b) Final address:

a.out :

main.o :

.text

e: R_X86_64_PLT32	swap-0x4	0000000000001181 <swap> (.text)
14: R_X86_64_PC32	buf	0000000000004010 <buf> (.data)

1a: R_X86_64_PC32	buf-0x4	0000000000004014 <buf> (.data)
23: R_X86_64_PC32	.rodata-0x4	0000000000002004 (.rodata)
2d: R_X86_64_PLT32	printf-0x4	0000000000001050 <printf@plt>

.eh_frame

20: R_X86_64_PC32	.text+0x0	0000000000001060
-------------------	-----------	------------------

swap.o :

.text

b: R_X86_64_PC32	buf	0000000000004014 <buf> (.data)
12: R_X86_64_PC32	bufp1-0x4	0000000000004030 <bufp1> (.data)
19: R_X86_64_PC32	bufp0-0x4	0000000000004020 <bufp0> (.data)
21: R_X86_64_PC32	.data-0x4	0000000000004018 <temp.1915>
28: R_X86_64_PC32	bufp1-0x4	0000000000004030 <bufp1> (.data)
2f: R_X86_64_PC32	bufp0-0x4	0000000000004020 <bufp0> (.data)
3a: R_X86_64_PC32	bufp1-0x4	0000000000004030 <bufp1> (.data)
40: R_X86_64_PC32	.data-0x4	0000000000004018 <temp.1915>

.data.rel

b: R_X86_64_PC32	buf	0000000000004014 <buf> (.data)
------------------	-----	--------------------------------

.eh_frame

20: R_X86_64_PC32	.text+0x0	0000000000001060
40: R_X86_64_PC32	.text+0x15	0000000000001075

other.o :

.text

1f: R_X86_64_PC32	buf-0x8	0000000000004010 <buf> (.data)
29: R_X86_64_PC32	buf-0x4	0000000000004014 <buf> (.data)

.eh_frame

20: R_X86_64_PC32	.text+0x0	0000000000001060
-------------------	-----------	------------------

Q2: Objdump 2

(a)

```
a) int main()
   401ce5:  f3 0f 1e fa      endbr64
   401ce9:  55               push  %rbp
   401cea:  48 89 e5         mov   %rsp,%rbp
b) int a=1, b=1;
   401ced:  c7 45 f8 01 00 00 00  movl  $0x1,-0x8(%rbp) -----(a = 1)
   401cf4:  c7 45 fc 01 00 00 00  movl  $0x1,-0x4(%rbp) -----(b = 1)

c) while(a<=10)
   401cfb:  eb 0e           jmp   401d0b <main+0x26>
d) b=b*a;
   401cfd:  8b 45 fc        mov   -0x4(%rbp),%eax
   401d00:  0f af 45 f8     imul  -0x8(%rbp),%eax
   401d04:  89 45 fc        mov   %eax,-0x4(%rbp)
e) a++;
   401d07:  83 45 f8 01     addl  $0x1,-0x8(%rbp)
f) while(a<=10)
   401d0b:  83 7d f8 0a     cmpl  $0xa,-0x8(%rbp)
   401d0f:  7e ec          jle   401cfd <main+0x18>
g) return b;
   401d11:  8b 45 fc        mov   -0x4(%rbp),%eax
   401d14:  5d             pop   %rbp
   401d15:  c3             retq
   401d16:  66 2e 0f 1f 84 00 00  nopw  %cs:0x0(%rax,%rax,1)
   401d1d:  00 00 00
```

(b) int a: -0x8(%rbp) , int b: -0x4(%rbp)

Q3: Objdump 3

(a)

```
a) struct data {  
    int sum;  
    int b[5];  
};  
int main()
```

```
401ce5:  f3 0f 1e fa      endbr64  
401ce9:  55               push  %rbp  
401cea:  48 89 e5         mov   %rsp,%rbp  
401ced:  48 83 ec 20      sub   $0x20,%rsp  
401cf1:  64 48 8b 04 25 28 00 mov  %fs:0x28,%rax  
401cf8:  00 00  
401cfa:  48 89 45 f8      mov   %rax,-0x8(%rbp)  
401cfe:  31 c0           xor   %eax,%eax
```

```
b) struct data rec1;  
   rec1.sum=0;
```

```
401d00:  c7 45 e0 00 00 00 00 movl  $0x0,-0x20(%rbp)
```

```
c) rec1.b[0]=2;  
   ⇒ ⇒
```

```
401d07:  c7 45 e4 02 00 00 00 movl  $0x2,-0x1c(%rbp)
```

```
d) rec1.sum=rec1.sum+rec1.b[0];
```

```
401d0e:  8b 55 e0         mov   -0x20(%rbp),%edx  
401d11:  8b 45 e4         mov   -0x1c(%rbp),%eax  
401d14:  01 d0           add   %edx,%eax  
401d16:  89 45 e0         mov   %eax,-0x20(%rbp)
```

```
e) return rec1.sum;
```

```
401d19:  8b 45 e0         mov   -0x20(%rbp),%eax  
401d1c:  48 8b 4d f8      mov   -0x8(%rbp),%rcx  
401d20:  64 48 33 0c 25 28 00 xor   %fs:0x28,%rcx  
401d27:  00 00  
401d29:  74 05           je    401d30 <main+0x4b>  
401d2b:  e8 d0 a0 04 00   callq 44be00 <__stack_chk_fail>  
401d30:  c9             leaveq  
401d31:  c3             retq  
401d32:  66 2e 0f 1f 84 00 00 nopw  %cs:0x0(%rax,%rax,1)  
401d39:  00 00 00  
401d3c:  0f 1f 40 00      nopl  0x0(%rax)
```

(b)

```
rec1.sum : -0x20(%rbp)  
rec1.b : -0x1c(%rbp)
```

Q4: Static Linking

Part (a)

Here, Module a1 has function main() and Module a2 has uninitialized global variable main.

Thus, function main in a1 is the strong symbol and main in a2 is weak symbol.

In such cases the linker always chooses the strong symbol. **Thus, the use of the symbol main in module a2 will resolve to the declaration of main in module a1.**

Checking the symbol table shows function main as the Global Symbol.

⇒ readelf -s a.out

⇒ ⇒ 1441: 0000000000401ce5 56 FUNC GLOBAL DEFAULT 7 main

0000000000401ce5 is the address of the main function.

Objdump also shows that references to main are for the main function.

The output produced shows same values. (Address of main function)

⇒ 0x401ce5

⇒ 0x401ce5

Part (b)

Here the linker gives error as both module b1 and module b2 have strong symbols.

b1 has function main() (Strong) and b2 has global initialized variable (int main = 1;) which is also strong. **This violates the rule that there should exist only one strong symbol with same name.** Thus the linker returns an **error**.

Part (c)

Here we get two symbols for main. Module c1.c has the Global Symbol of main function and c2.c has Local Symbol static int main = 1.

⇒ readelf -s a.out

⇒ ⇒ 167: 00000000004be0f0 4 OBJECT LOCAL DEFAULT 21 main

⇒ ⇒ 1442: 0000000000401ce5 56 FUNC GLOBAL DEFAULT 7 main

The printf in c1.c can only reference the strong symbol (main function) because static variables and functions can be accessed only from their file thus the address printed is that of main function. Whereas in c2.c we have a local symbol (static int main) thus according to the scoping rules the printf statement prints the address of Local Symbol main.

Thus the output produced shows different addresses:

⇒ 0x401ce5

⇒ 0x4be0f0

The use of the symbol main in module c1 will resolve to the declaration of main in module c1 & the use of the symbol main in module c2 will resolve to the declaration of main in module c2.

Q5: Memory Layout

Symbol	Section	Region Number
main()	.text (Read Only)	5
f()	.text (Read Only)	5
malloc()	.plt.sec (Read Only)	5
printf()	.plt.sec (Read Only)	5
x	.data (Read/Write)	4
y	.data (Read/Write)	4
a	.bss (Read/Write)	4
k	.bss (Read/Write)	4
p	User Stack	1

*** p is not a symbol ??? ***

Q6: Link Issue!

- a) Output :
⇒ 5 13
⇒ 13 5

Here we can see that printf in f() outputs swapped values of x and y, whereas printf in main() prints the correct values of x and y.

rec in module a1.c is Strong Symbol and that in module a2.c is Weak Symbol. Thus, the linker chooses struct rec of module a1.c.

We need to notice that struct in a2.c has y before and then x. Thus, references made to x in a2.c have address corresponding to that of y in struct rec in module a1.c.

The assembly code generated for a2.c mention location for rec.x as rec+0x4 and for rec.y as rec+0x0. But after linking rec refers to definition in module a1.c. Thus, rec+0x0

points to x and rec+0x4 points to y. Thus, the output produced is according to addresses of the assembly generated.

main:

```
printf(" %d %d\n", rec.x, rec.y);
115b: 8b 15 b3 2e 00 00    mov  0x2eb3(%rip),%edx    # 4014 <rec+0x4>
1161: 8b 05 a9 2e 00 00    mov  0x2ea9(%rip),%eax    # 4010 <rec>
1167: 89 c6                mov  %eax,%esi
1169: 48 8d 3d 94 0e 00 00 lea  0xe94(%rip),%rdi     # 2004
<_IO_stdin_used+0x4>
1170: b8 00 00 00 00      mov  $0x0,%eax
1175: e8 d6 fe ff         callq 1050 <printf@plt>
```

f1:

```
printf(" %d %d\n", rec.x, rec.y);
1189: 8b 15 81 2e 00 00    mov  0x2e81(%rip),%edx    # 4010 <rec>
118f: 8b 05 7f 2e 00 00    mov  0x2e7f(%rip),%eax    # 4014 <rec+0x4>
1195: 89 c6                mov  %eax,%esi
1197: 48 8d 3d 6e 0e 00 00 lea  0xe6e(%rip),%rdi     # 200c
<_IO_stdin_used+0xc>
119e: b8 00 00 00 00      mov  $0x0,%eax
11a3: e8 a8 fe ff         callq 1050 <printf@plt>
```

We can see in objdump that the values in eax and edx are swapped for main() and f1().

-----That implies that referring to rec1.x in main() function points to location 0000000000004014 and rec1.y points to 0000000000004010.

Whereas referring to rec1.x in f1() function points to location 0000000000004010 and rec1.y in f1() function points to 0000000000004014.-----

b) Locations that need relocation:

test1.c

.text

1. e: R_X86_64_PLT32 f1-0x4
2. 14: R_X86_64_PC32 rec
3. 1a: R_X86_64_PC32 rec-0x4
4. 23: R_X86_64_PC32 .rodata-0x4
5. 2d: R_X86_64_PLT32 printf-0x4

.eh_frame

1. 20: R_X86_64_PC32 .text+0x0

test2.c

.text

1. a: R_X86_64_PC32 rec-0x4
2. 10: R_X86_64_PC32 rec

3. 19: R_X86_64_PC32 .rodata-0x4
 4. 23: R_X86_64_PLT32 printf-0x4
 .eh_frame
 1. 20: R_X86_64_PC32 .text+0x0

c) Final addresses:

.text

main():

e: R_X86_64_PLT32	f1-0x4	0000000000001181 <f1>
14: R_X86_64_PC32	rec	0000000000004014 <rec+0x4>
1a: R_X86_64_PC32	rec-0x4	0000000000004010 <rec>
23: R_X86_64_PC32	.rodata-0x4	0000000000002004 <_IO_stdin_used+0x4>
2d: R_X86_64_PLT32	printf-0x4	1050 <printf@plt>

f1():

a: R_X86_64_PC32	rec-0x4	0000000000004010 <rec>
10: R_X86_64_PC32	rec	0000000000004014 <rec+0x4>
19: R_X86_64_PC32	.rodata-0x4	000000000000200c <_IO_stdin_used+0xc>
23: R_X86_64_PLT32	printf-0x4	1050 <printf@plt>

.eh_frame

20: R_X86_64_PC32	.text+0x0	0000000000001060
-------------------	-----------	------------------

Q7: Static Linking - 2

- a) We get error in line in module f2.c:

⇒ int z = a[3];

This is because int a[] is defined as an extern variable in f2.c. Thus, it can not be used in that module.

- b) The output observed for (5 == fn()) is 0. That is because d is no more 5 after statement: c = 100; in module f1.c

In f1.c the assembly code generated is according to c being a double variable (8 bytes). But { int c = 58; } is the Strong Symbol. Thus the reference is made to this. Variable d occupies address right after c. Thus, when we execute c = 100 in f1.c, 8 bytes of data is written overriding d's actual value. Thus, d is not 5 anymore making e != 5. fn() doesn't return 5 and thus the output is 0.

- c) Locations that need relocation:

Address of locations that need relocation are mentioned with address relative to their section in the .o file.

f1.c

.text

1. 11: R_X86_64_PC32 .rodata+0x4
2. 19: R_X86_64_PC32 c-0x4
3. 20: R_X86_64_PC32 c-0x4
4. 2a: R_X86_64_PLT32 fn-0x4
5. 41: R_X86_64_PC32 .rodata-0x4
6. 4b: R_X86_64_PLT32 printf-0x4

.data.rel.local -----doubt

1. 0: R_X86_64_64 a+0xc

.data.rel -----doubt

1. 0: R_X86_64_64 b+0x0

f2.c

.text

1. 0e: R_X86_64_PC32 d-0x4
2. 1d: R_X86_64_PC32 .rodata-0x4
3. 27: R_X86_64_PLT32 printf-0x4

Final Addresses

.text

main()

Locations that need relocation	Final address
11: R_X86_64_PC32 .rodata+0x4	0000000000494010 <_IO_stdin_used+0x10>
19: R_X86_64_PC32 c-0x4	00000000004be118 <c>

20: R_X86_64_PC32	c-0x4	00000000004be118 <c>
2a: R_X86_64_PLT32	fn-0x4	0000000000401d40 <fn>
41: R_X86_64_PC32	.rodata-0x4	0000000000494008 <_IO_stdin_used+0x8>
4b: R_X86_64_PLT32	printf-0x4	0000000000410d40 <_IO_printf>

fn()

0e: R_X86_64_PC32	d-0x4	00000000004be11c <d>
1d: R_X86_64_PC32	.rodata-0x4	0000000000494018 <_IO_stdin_used+0x18>
27: R_X86_64_PLT32	printf-0x4	0000000000410d40 <_IO_printf>

.data.rel.local

0: R_X86_64_64	a+0xc	00000000004be0f0 <a> (.data)
----------------	-------	------------------------------

.data.rel.

0: R_X86_64_64	b+0x0	00000000004be120 (.data)
----------------	-------	------------------------------

- a. Symbols with relocation type R_X86_64_PC32 are relative to program counter.
 1. c
 1. d
- b. Symbols with relocation type R_X86_64_PLT32 use the Program linkage table.

“In the case of a position-independent code, the call to a function be done through a table called the Procedure Linkage Table which is used by the dynamic linker at runtime. This type has been chosen in case we would link the nothing.o in a shared library and link the executable with this dynamic library. In the case all is statically linked, the linker will consider it will have to do the same job as if the relocation type was R_X86_64_PC32 relocation.”

Source:

<https://stac47.github.io/c/relocation/elf/tutorial/2018/03/01/understanding-relocation-elf.html>

Since, we assume static linking, relocation of the following symbols is relative to program counter.

1. printf
2. fn

- c. Symbols with relocation type R_X86_64_64 are PC independent
 - 1. a
 - 2. b