

# CS685: DATA MINING ENSEMBLE METHODS

Arnab Bhattacharya  
arnabb@cse.iitk.ac.in

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur  
<http://web.cse.iitk.ac.in/~cs685/>

1<sup>st</sup> semester, 2021-22  
Mon 1030-1200 (online)

# Ensemble Methods

- Combine *multiple* classifiers to increase classification accuracy
- A single instance of a classifier may do a mistake, but it is unlikely that many of them will
- **Ensemble** methods are also called **classifier combination** methods

# Ensemble Methods

- Combine *multiple* classifiers to increase classification accuracy
- A single instance of a classifier may do a mistake, but it is unlikely that many of them will
- **Ensemble** methods are also called **classifier combination** methods
- Final class is determined by simply using *majority voting* of outputs of individual classifiers
- Voting may be *weighted* as well

# Ensemble Methods

- Combine *multiple* classifiers to increase classification accuracy
- A single instance of a classifier may do a mistake, but it is unlikely that many of them will
- **Ensemble** methods are also called **classifier combination** methods
- Final class is determined by simply using *majority voting* of outputs of individual classifiers
- Voting may be *weighted* as well
- Classifiers of same or different types can be combined
- Same type: **bagging**, **boosting**
- Different type: **stacking**

# Ensemble Methods

- Combine *multiple* classifiers to increase classification accuracy
- A single instance of a classifier may do a mistake, but it is unlikely that many of them will
- **Ensemble** methods are also called **classifier combination** methods
- Final class is determined by simply using *majority voting* of outputs of individual classifiers
- Voting may be *weighted* as well
- Classifiers of same or different types can be combined
- Same type: **bagging**, **boosting**
- Different type: **stacking**
- Assumes *no* class imbalance problem, i.e., classes are almost equally represented

# Bagging

- Short form of **bootstrap aggregating**
- Suppose, initial training set size is  $d$
- Create another training set of *same* size  $d$  by *sampling with replacement*
  - Equivalent to deleting some objects and replicating some others
- Repeat  $t$  times
- Build  $t$  classifiers of the *same* type on each of the training samples
- Use majority voting on the  $t$  classifiers for final classification

# Bagging

- Short form of **bootstrap aggregating**
- Suppose, initial training set size is  $d$
- Create another training set of *same* size  $d$  by *sampling with replacement*
  - Equivalent to deleting some objects and replicating some others
- Repeat  $t$  times
- Build  $t$  classifiers of the *same* type on each of the training samples
- Use majority voting on the  $t$  classifiers for final classification
- Random forests can be considered as a type of bagging

## 0.632 Bootstrap



## 0.632 Bootstrap

- Creating a sample of size  $d$  from a set of size  $d$  by *sampling with replacement*

## 0.632 Bootstrap

- Creating a sample of size  $d$  from a set of size  $d$  by *sampling with replacement*
- Probability of a particular object being chosen in a try is  $1/d$
- Therefore, probability of not being chosen is  $1 - 1/d$

## 0.632 Bootstrap

- Creating a sample of size  $d$  from a set of size  $d$  by *sampling with replacement*
- Probability of a particular object being chosen in a try is  $1/d$
- Therefore, probability of not being chosen is  $1 - 1/d$
- Thus, probability of *not* being chosen at all in any of the  $d$  tries is  $(1 - 1/d)^d$

## 0.632 Bootstrap

- Creating a sample of size  $d$  from a set of size  $d$  by *sampling with replacement*
- Probability of a particular object being chosen in a try is  $1/d$
- Therefore, probability of not being chosen is  $1 - 1/d$
- Thus, probability of *not* being chosen at all in any of the  $d$  tries is  $(1 - 1/d)^d$
- For large  $d$ ,  $(1 - 1/d)^d \rightarrow 1/e = 0.368$
- Hence, probability that an object is chosen at least once, i.e., it is represented in the sample is  $1 - 0.368 = 0.632$
- Hence, the method is called **0.632 bootstrap**

# Boosting

- Classifiers of the same type
- Uses majority voting of classifiers
- May use weights
- Models are *not* built independently
- Each model is built *successively* on the previous ones

# Boosting

- Classifiers of the same type
- Uses majority voting of classifiers
- May use weights
- Models are *not* built independently
- Each model is built *successively* on the previous ones
- Each new model concentrates more on the training objects mis-classified by previous models
- Each training set is of the *same* size  $d$  and is constructed using *sampling with replacement*

# Adaboost

- A popular boosting algorithm is **Adaboost**
- Each training object is weighted (initially with the same weight)
- For a model, the *error* in classification is computed as the ratio of the sum of weights of mis-classified objects to total weight
- Total weight is normalized to 1
- Hence, error  $e$  is

$$e = \sum_{p \text{ is mis-classified}} w(p)$$

# Adaboost

- A popular boosting algorithm is **Adaboost**
- Each training object is weighted (initially with the same weight)
- For a model, the *error* in classification is computed as the ratio of the sum of weights of mis-classified objects to total weight
- Total weight is normalized to 1
- Hence, error  $e$  is

$$e = \sum_{p \text{ is mis-classified}} w(p)$$

- Models are successively built by *modifying* the weights of training objects
- If a training object is *correctly* classified, its weight is *decreased*
- Consequently, weights of mis-classified objects increase



# Adaboost

- A popular boosting algorithm is **Adaboost**
- Each training object is weighted (initially with the same weight)
- For a model, the *error* in classification is computed as the ratio of the sum of weights of mis-classified objects to total weight
- Total weight is normalized to 1
- Hence, error  $e$  is

$$e = \sum_{p \text{ is mis-classified}} w(p)$$

- Models are successively built by *modifying* the weights of training objects
- If a training object is *correctly* classified, its weight is *decreased*
- Consequently, weights of mis-classified objects increase
- For the next model, sampling with replacement is done using the new weights as probabilities

# Updating the Weights

- Initially, the weights of training objects are same
- Thus, model 1 is built by simple sampling with replacement from training set
- Weights of all *correctly* classified objects are modified using the error  $e$  of the model

$$w(q) = w(q) \times \frac{e}{(1 - e)} \text{ where } q \text{ is correctly classified}$$

# Updating the Weights

- Initially, the weights of training objects are same
- Thus, model 1 is built by simple sampling with replacement from training set
- Weights of all *correctly* classified objects are modified using the error  $e$  of the model

$$w(q) = w(q) \times \frac{e}{(1 - e)} \text{ where } q \text{ is correctly classified}$$

- If  $e = 0$ , then all objects are correctly classified and model building is stopped
- If  $e > 0.5$ , then more than half the objects are mis-classified
- Such a model is deemed useless and is discarded

# Updating the Weights

- Initially, the weights of training objects are same
- Thus, model 1 is built by simple sampling with replacement from training set
- Weights of all *correctly* classified objects are modified using the error  $e$  of the model

$$w(q) = w(q) \times \frac{e}{(1 - e)} \text{ where } q \text{ is correctly classified}$$

- If  $e = 0$ , then all objects are correctly classified and model building is stopped
- If  $e > 0.5$ , then more than half the objects are mis-classified
- Such a model is deemed useless and is discarded
- Sampling is repeated till a suitable stopping criterion

# Updating the Weights

- Initially, the weights of training objects are same
- Thus, model 1 is built by simple sampling with replacement from training set
- Weights of all *correctly* classified objects are modified using the error  $e$  of the model

$$w(q) = w(q) \times \frac{e}{(1 - e)} \text{ where } q \text{ is correctly classified}$$

- If  $e = 0$ , then all objects are correctly classified and model building is stopped
- If  $e > 0.5$ , then more than half the objects are mis-classified
- Such a model is deemed useless and is discarded
- Sampling is repeated till a suitable stopping criterion
- Final models are weighted for majority voting according to  $\log \frac{1-e_i}{e_i}$  where  $e_i$  is the error of model  $i$

# Discussion

- Models are necessarily of the same type
- Model outputs are assumed to be independent for voting

# Discussion

- Models are necessarily of the same type
- Model outputs are assumed to be independent for voting
- Bagging is simple
- Boosting progressively concentrates on harder training objects
- Boosting may overfit

- Models are necessarily of the same type
- Model outputs are assumed to be independent for voting
- Bagging is simple
- Boosting progressively concentrates on harder training objects
- Boosting may overfit
- Assumes no *class imbalance*
- If a class is under-represented, it has to be *over-sampled*
- If a class is over-represented, it has to be *under-sampled*
- May need to take into account the relative populations while computing the mis-classification error



# Stacking

- **Stacking** or **stacked generalization** combines models of *different* types

# Stacking

- **Stacking** or **stacked generalization** combines models of *different* types
- Individual classifiers are models of **level 0**
- Builds another classifier of **level 1** that combines the *outputs* of level-0 classifiers
- Level-1 classifier uses *only* the outputs of level-0 models, i.e., just their class labels and *not* the attributes of the training objects
- May use errors as well

# Stacking

- **Stacking** or **stacked generalization** combines models of *different* types
- Individual classifiers are models of **level 0**
- Builds another classifier of **level 1** that combines the *outputs* of level-0 classifiers
- Level-1 classifier uses *only* the outputs of level-0 models, i.e., just their class labels and *not* the attributes of the training objects
- May use errors as well
- Level-1 classifier is thus just an arbiter
- It should be simple

# Stacking

- For better results, training sample is divided into two parts  $D_0$  and  $D_1$
- Level-0 classifiers are trained using *only*  $D_0$  and *not*  $D_1$
- Level-1 classifier is trained using *only*  $D_1$  and *not*  $D_0$ 
  - $D_1$  training objects are simply passed through level-0 classifiers to get their outputs
  - These outputs act as inputs to the level-1 classifier
  - $D_1$  is not used for *training* level-0 classifiers at all

